

The Co-dfns Compiler

Aaron W. Hsu

July 10, 2022

Co-dfns Compiler: High-performance, Parallel APL Compiler
Copyright © 2011-2022 Aaron W. Hsu <arcfide@sacrideo.us>

This program is free software: you can redistribute it and/or modify it under the terms of the GNU Affero General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Affero General Public License for more details.

You should have received a copy of the GNU Affero General Public License along with this program. If not, see <http://gnu.org/licenses>.

This program is available under other license terms. Please contact Aaron W. Hsu <arcfide@sacrideo.us> for more information.

Contents

1	Introduction	6
1.1	How to Read a WEB	6
2	User’s Guide	6
3	Co-dfns Architecture	6
3.1	Global Settings	10
3.2	The Fix API	13
3.3	The User Command API	15
3.4	AST Record Structure	15
3.5	Converters between parent and depth vectors	15
4	Testing	16
5	Co-dfns Compiler	17
5.1	Parser	17
5.1.1	Output of the Parser	18
5.1.2	Handling Parsing Errors	20
5.1.3	Tokenizing the Input	21
5.1.4	Parsing Token Stream	22
5.2	Compiler Transformations	23
5.3	Code Generator	24
5.4	Backend C Compiler Interface	26
5.5	Linking with Dyalog	27
5.6	Runtime	31
5.6.1	GPU C Runtime Kernel	33
6	Language Features	47
6.1	Valid source input character set	47
6.2	Comments and Whitespace	48
6.3	Numbers	48
6.4	Strings and characters	49
6.5	Variables	49
6.6	Arrays	50
6.7	Primitives	50
6.7.1	APL Primitives	51
6.7.2	System Functions and Variables	51
6.8	Brackets	52
6.8.1	Indexing	52
6.8.2	Axis Operator	53
6.9	Bindings and Types	54
6.10	Assignments	55
6.11	Expressions	58
6.11.1	Value Expressions	59

6.11.2	Function Expressions	60
6.12	Trains	60
6.13	Functions	61
6.13.1	D-fns	61
6.13.2	Trad-fns	63
6.14	Guards	64
6.14.1	Error Guards	64
6.15	Labels	64
6.16	Statements	65
6.16.1	What is a keyword?	65
6.16.2	Namespaces	65
6.16.3	Structured Programming Statements	68
7	Runtime Primitives	68
7.1	Addition/Identity	68
7.2	And (Logical)	69
7.3	Bracket	69
7.4	Catenate (First/Last Axis)	69
7.5	Circle/Trigonometrics	69
7.6	Commute	69
7.7	Compose	69
7.8	Convolve	69
7.9	Decode	70
7.10	Disclose	70
7.11	Division/Reciprocal	70
7.12	Drop	70
7.13	Each	70
7.14	Enclose	70
7.15	Encode	70
7.16	Equal	71
7.17	Exponent	71
7.18	Factorial/Binomial	71
7.19	Fast Fourier Transforms	71
7.20	Find	71
7.21	Grade Down	71
7.22	Grade Up	72
7.23	Greater Than	72
7.24	Greater Than or Equal	72
7.25	Index	72
7.26	Index Generator	72
7.27	Inner Product	72
7.28	Intersection	72
7.29	Left	73
7.30	Less Than	73
7.31	Less Than or Equal	73
7.32	Logarithm	73

7.33 Match	73
7.34 Matrix Division	73
7.35 Maximum/Ceiling	73
7.36 Membership	74
7.37 Minimum/Floor	74
7.38 Multiplication	74
7.39 Nest/Partition	74
7.40 Not	74
7.41 Not And (Logical)	74
7.42 Not Equal	74
7.43 Not Match	75
7.44 Not Or (Logical)	75
7.45 Or (Logical)	75
7.46 Outer Product	75
7.47 Power	75
7.48 Rank	75
7.49 Reduce	75
7.50 Roll	76
7.51 Rotate (First/Last Axis)	76
7.52 Residue	76
7.53 Right	76
7.54 Scalar Each	76
7.55 Scan	77
7.56 Shape	77
7.57 Subtraction	77
7.58 Take	77
7.59 Transpose	77
7.60 Union	77
8 Utilities	78
8.1 Must haves	78
8.2 AST Pretty-printing	79
8.3 Debugging utilities	80
8.4 Reading and Writing Files	81
8.5 XML Rendering	82
8.6 Detecting the Operating System	82
9 Developer Infrastructure	83
9.1 Building the Compiler	83
9.1.1 Tangling the Source	83
9.1.2 Weaving the Source	85
9.2 Building the Runtime	87
9.3 Loading the Compiler	90

July 10, 2022

`codfns.nw` 5

10 Index	91
10.1 Chunks	91
10.2 Identifiers	93
11 GNU AFFERO GPL	96

1 Introduction

1.1 How to Read a WEB

2 User's Guide

3 Co-dfns Architecture

This section describes the “big picture” parts of the Co-*dfns* compiler. The intent here is to try to show how all of the various moving parts of the compiler fit together, to provide a sort of road map that will give you a precise plan for understanding how the various components affect one another. One of the most important things to understand in any compiler is the net effect a local change in the code can have on the rest of the system, so I hope that this section will help to clarify this.

The design of the Co-*dfns* compiler is one of austerity and minimalism. My intent is, was, and hopefully shall remain that of producing an exceptionally clear design that avoids or eliminates unnecessary code and complexity within the design. I attack this problem in many ways, but I primarily attempt to do this by both reducing the size of the code surface in total, that is, write less code, as well as reducing the number of entry points and paths through that code. In other words, my ideal design is one in which you enter the compiler in some limited, but well defined and useful set of entry points, and then proceed in a linear fashion through the code as the execution path, resulting finally in your result. This is the “ultimate” in data flow, functionally oriented programming.

The ramifications of this design choice implies a few important things. Firstly, it implies that I reduce and eliminate any code that represents boilerplate or that does not actively contribute to the “big picture” of the code. This is required in an extreme degree if I am to reduce the overall complexity of the design. This also implies that there is very little intentional redundancy in the shape and style of the source, making it very terse and compact. Since there are intentionally very few entry and exit points through the control flow of the code, this reduces the number of dependencies for me to be aware of when dealing with a single piece of code, but this also comes at the cost of not being able to see many examples of the interfaces with that code. Often, there will be one, and only one place, in which a given piece of code is used, and I do not want the code to needlessly store excess information in its source that doesn't need to be there.

This all culminates in something that can be quite shocking at first: making a change to the source is almost always a big deal. If

all the source code is meaningful and carefully constructed, this also means that changing this code is almost always non-trivial, because if the code represented something trivial, I would have tried to remove it from the code so that only the “big things” were in the code itself. Thus, anyone who wishes to view and read the compiler code should take it upon themselves to appreciate the way in which the code flows together, and how the flow of the program runs, as doing so will be essential to understanding how to make changes to the source without breaking something. Fortunately, this does come with the intended benefits of a very short and simple codebase that has clear flow through the system, it just means that if you want to change something, make sure you realize that you are almost always likely to be working at the “architectural” level, rather than at the small and trivial level of details.

The compiler is designed to fit into a single Dyalog APL namespace, and importantly, we do not define additional nested namespaces or other forms of name hiding. I intentionally want to restrict the namespace to a single global one. This single global namespace should therefore contain the carefully curated names that matter, and any that do not matter should, ideally, not be defined or used. The namespace itself can be divided into three main groupings: the public facing entry-points into the system, the compiler logic itself, and the utilities or other elements that serve to support the others. This gives use the following code outline.

```
7  ⟨*7⟩≡
    :Namespace codfns

    ⟨Global Settings 10a⟩
    ⟨The Fix API 13⟩
    ⟨User-command API 15a⟩

    ⟨Parser 17⟩
    ⟨Compiler 23⟩
    ⟨Code Generator 25b⟩
    ⟨Interface to the backend C compiler 26⟩
    ⟨Linking with Dyalog 27⟩

    ⟨Must Have APL Utilities 78⟩
    ⟨Basic tie and put utilities 81c⟩
    ⟨The opsys utility 82b⟩
    ⟨AST Record Structure 15b⟩
    ⟨Converters between parent and depth vectors 15c⟩
    ⟨XML Rendering 82a⟩
    ⟨Pretty-printing AST trees 79⟩
```

:EndNamespace

Root chunk (not used in this document).

Defines:

`codfns`, used in chunks 8, 16b, 24d, 26, 32b, 34a, 41, 47b, 58, 68, 81b, and 83–89.

This $\langle * 7 \rangle$ chunk is meant to be stored to a file. We have a build system for doing this that depends on the contents of the $\langle \textit{Tangle Commands 8} \rangle$ chunk. Thus, we follow the convention here of updating the contents of the $\langle \textit{Tangle Commands 8} \rangle$ chunk each time that we initially define a new chunk that is intended to be output to a file during the tangling process. See more about the build infrastructure later in this document.

8 $\langle \textit{Tangle Commands 8} \rangle \equiv$
`echo "Tangling codfns.apln..."`
`notangle codfns.nw > src/codfns.apln`

This definition is continued in chunks 16b, 32b, 34a, 41b, 47b, 58b, 68b, 81b, 84, 86, and 87c.

This code is used in chunk 83.

Defines:

`codfns.apln`, never used.

Uses `codfns 7` and `src 88`.

The primary user-facing interfaces into the compiler are *⟨The Fix API 13⟩* and the *⟨User-command API 15a⟩*. These are the ways that you primarily drive the entire compiler. I intentionally expose the rest of the compiler interfaces without hiding them so that people who wish to leverage these other parts of the system without using the “entire” compiler pipeline are able to do so, but I do not consider this a public interface.

This distinction matters because of our testing philosophy and our version numbering. Generally speaking, our version numbering scheme only tracks a major or minor change in the compiler when the externally facing interfaces receive some fundamental changes. Changes to the internal changes are *not* considered for this versioning scheme. Moreover, since I intend for there to be great freedom in changing and altering the behavior of these internal pipeline interfaces, these interfaces are not directly tested, and the test suite should *not* include testing against these internal interfaces. We philosophically only test against the external interfaces, and eschew internal unit tests.¹

The utility functions defined below the core compiler pipeline represent functionality that is tangential to the main compiler operation. However, these utilities also tend to represent some specific insight into the design of the compiler. Understanding the core AST structure and design as well as getting a grip on how to manipulate the core tree manipulation structures are vital to understanding the rest of the code. Therefore, this section spends more time on discussing these topics before the upcoming sections dealing with a more detailed exposition of the compiler itself. However, there are utilities that we consider more advanced, such as the pretty-printing functions and XML rendering that are topics of interest to advanced users of the compiler, but which are not part of the main compiler pipeline. Even though these functions have intentionally general application and are likely to be useful not only to those working on the compiler itself but also to those who are using more advanced compiler features, these utilities are not critical to a deep understanding of the compiler, so these are not discussed in this section. Instead, we discuss those topics in the section on developer tooling and infrastructure concerns.

The remaining parts of this section will describe the external facing interfaces to the compiler as well as the core underlying data structures and idioms that form the underlying skeleton and foundation for writing and working with any aspect of the compiler. These are all feature and component agnostic elements of the system that do not belong solely to only a single part, but that impact all other

¹You can read more of my opinions on this matter in my article, “The Fallacy of Unit Testing”.

elements of the compiler source code, and so it pays especially well to pay attention and understand this code to a high degree.

3.1 Global Settings

There are some global options that we assume to exist throughout the compiler. These set the standard behaviors as well as serve as knobs that can be tweaked in some cases to identify what behaviors we want from the rest of the compiler.

First, we have a set of read-only global constants that are defined to configure our APL environment. These are the typical ones, and we try to stick to the defaults, except that we are sane, and thus we use `⎕IO` set to 0.

10a $\langle \textit{Global Settings } 10a \rangle \equiv$
`⎕IO ⎕ML ⎕WX←0 1 3`

This definition is continued in chunks 10–12.

This code is used in chunk 7.

Defines:

`⎕IO`, used in chunk 80.

`⎕ML`, used in chunk 80.

`⎕WX`, never used.

Additionally, we set a `VERSION` constant to track changes to the system through the distributions. We use semantic versioning² as our versioning scheme. That being said, we also do not have particular qualms about changing the public API at a rapid pace, provided that we document this.

10b $\langle \textit{Global Settings } 10a \rangle + \equiv$
`VERSION←4 1 0`

This code is used in chunk 7.

Defines:

`VERSION`, never used.

²<https://semver.org/>

We depend on ArrayFire³ for much of our GPU backend functionality. This means we need to know two things, where ArrayFire is installed and which ArrayFire backend we should use when compiling. We only really need to know where ArrayFire is installed on UNIX style systems, as these systems seem to be much more variable in this regard, and there is an environment variable that we can use in Windows to find out where ArrayFire is installed more conveniently on that platform. We default to using 'cuda' as our main option, but we also support the following options for `AF_LIB`:

```
cuda opengl cpu
```

Using '' for `AF_LIB` will use ArrayFire's unified backend, but we don't default to this because we have seen some issues on some platforms with reliability problems. To avoid this, we choose to use `cuda` as the default, which tends to either work or fail explicitly, which allows the user to respond rather than crashing ungracefully in the case of the unified backend.

The least reliable backend we have seen is the `opengl` one, which seems to be more hit or miss depending on the underlying stability of the OpenCL drivers that are installed on the user's system. In particular, some Linux OpenCL installations seem to be particularly fragile. In such cases, always make sure that a good, solid OpenCL library is being used.

```
11 <Global Settings 10a>+≡
    AFΔPREFIX←'/opt/arrayfire'
    AFΔLIB←'cuda'
```

This code is used in chunk 7.

Defines:

`AFΔLIB`, used in chunks 15a, 26, and 89a.
`AFΔPREFIX`, used in chunk 26.

³<https://arrayfire.com/>

On Windows, we rely on the Visual Studio C/C++ compiler to build our runtime and user code. We have settled on trying to stay as up to date with this as possible. However, there are many different installation paths used by Visual Studio, which can make it difficult to know where to look unless we hardcode each location. Instead, we assume that Visual Studio will not be a primary interest to our users, making it likely that they will be installing Visual Studio only as a dependency for using Co-*dfns*. In this case, it is likely that they will be using the Community version. Thus, we default to using the latest version of Visual Studio of which we are aware and using the Community version of this, which Microsoft does not charge for.

If a different version of Visual Studio is installed, then it is important to figure out what the right path should be to locate the Visual Studio installation. The main thing we need to get from this path is access to the `vcvarsall.bat` batch file. This file configures the `cmd.exe` environment to be able to find the Visual Studio compiler and work in the right way. In the 2002 Community addition, and apparently most new versions of Visual Studio, this is located in the `VC\Auxiliary\Build\` subdirectory of the main installation folder. When changing this path, we want to make sure that the following path points to the correct `vcvarsall.bat` file:

```
VSΔPATH, '\VC\Auxiliary\Build\vcvarsall.bat'
```

Most users will simply need to alter `Community` to match the edition of Visual Studio 2022 that they have installed on their system.

```
12 <Global Settings 10a>+≡
    VSΔPATH←'\Program Files\Microsoft Visual Studio'
    VSΔPATH,←'\2022\Community'
```

This code is used in chunk 7.

Defines:

VSΔPATH, used in chunks 26 and 89a.

3.2 The Fix API

One of the core entry points into the compiler is through the `Fix` function. This function is designed to mimic and more or less replace the use of the `FIX` function found in Dyalog APL. Its design models that behavior, and it is important as an entry-point because it exercises most of the core elements of the compiler. In particular, the design of the compiler’s pipeline is demonstrated most fully in this function.

Parse → Compile → Generate → Backend → Link

The interfaces to the `FIX` function and the Co-dfns `Fix` function differ in a few key ways. The left argument to `Fix` is a character vector giving the name to use when generating files and other artifacts. This does *not* affect the name of the resulting namespace, since that is defined, if at all, in the file source itself. The α argument only affects the name of the files and other outputs that `Fix` generates.

We also print out which part of the compiler we are in when we enter that “phase”. Doing this helps to give us an intuitive sense of how fast each phase is and whether one phase is taking an abnormally long time or not. It also helps in debugging.

```
13  ⍎The Fix API 13≡
    Fix←{
      _←a n s src←PS ω→_←'P'
      _←          TT _←_←'C'
      _←          GC _←_←'G'
      _←          α CC _←_←'B'
      n NS _←_←'L'
    }
```

This code is used in chunk 7.

Defines:

`Fix`, used in chunk 15a.

Uses `PS` 17 and `src` 88.

The input requirements for `Fix` are not listed in the definition itself, because both the parser `PS` and the `Fix` function need to use the same basic checks, and since the `Fix` function calls the parser as its first entry point, it doesn't make much sense to duplicate that work in both places. The requirements are as follows:

- Scalar/Vector
- Character type
- Simple or Vector of Vectors

We generate a `DOMAIN ERROR` if the inputs are not well-formed.

```
14a  <Verify source input ω, set IN 14a>≡
      IN←ω

      err←'PARSER EXPECTS SCALAR OR VECTOR INPUT'
      1<≠pIN:err □SIGNAL 11

      err←'PARSER EXPECTS SIMPLE OR VECTOR OF VECTOR INPUT'
      2<|≡IN:err □SIGNAL 11

      <Normalize the input formatting 14b>

      err←'PARSER EXPECTS CHARACTER ARRAY'
      0≠10|□DR IN:err □SIGNAL 11
```

This code is used in chunk 17.
Uses `SIGNAL 20b`.

The input formatting that is accepted means that newlines could be denoted either with `LF`, `CR`, or `CRLF` sequences inside of the vectors themselves or they could be denoted by having separate vectors for the various lines, or even a mixture of both. To simplify this situation we want to normalize them so that we are always dealing with some combination of `LF`, `CR`, and `CRLF` sequences within the file itself, rather than dealing with the nested situation. This ensures that after verification of the input, everything will work off of the same format. We intentionally put a newline at the end of the file even if we may not require one because it is possible that we are dealing with a file that is missing its final newline. By always adding one, we ensure that every line in the input is always terminated by a line ending. Life is also simpler if we just use `LF` as our line ending instead of something else, this means that future code must be aware that there could be mixed line endings in the file.

```
14b  <Normalize the input formatting 14b>≡
      IN←ε(⊆IN), ``□UCS 10
```

This code is used in chunk 14a.

3.3 The User Command API

15a $\langle \text{User-command API } 15a \rangle \equiv$

```

  ▽ Z ← Help _
  Z ← 'Usage: <object> <target> [-af={cpu,opencl,cuda}]'
  ▽

  ▽ r ← List
  r ← NS''1p<Θ ◇ r.Name ←, ``c'Compile' ◇ r.Group ← c'CODFNS'
  r[0].Desc ← 'Compile an object using Co-dfns'
  r.Parse ← c'2S -af=cpu opencl cuda '
  ▽

  ▽ Run(C I); Convert; in; out
  A Parameters
  A      AFΔLIB      ArrayFire backend to use
  Convert ← {α(□SE.SALT.Load'[SALT]/lib/NStoScript -noname').ntgennscode ω}
  in out ← I.Arguments ◇ AFΔLIB ← I.af'' > ∼ I.af ≡ 0
  S ← (c':Namespace ', out), 2 ↓ 0 0 0 out Convert ##.THIS.⊕ in
  → 0 / ∼ 'Compile' ≠ C
  {##.THIS.⊕ out, '← ω'} out Fix S → □ EX'##.THIS.', out
  ▽

```

This code is used in chunk 7.
Uses AFΔLIB 11 and Fix 13.

3.4 AST Record Structure

15b $\langle \text{AST Record Structure } 15b \rangle \equiv$

```

  fΔ ← 'ptknfsrdx'
  NΔ ← 'ABCEFGKLMNOPSVZ'
  A B C E F G K L M N O P S V Z ← 1 + ι 15

```

This code is used in chunk 7.

3.5 Converters between parent and depth vectors

15c $\langle \text{Converters between parent and depth vectors } 15c \rangle \equiv$

```

  P2D ← {z ← ∼ ι ≠ ω ◇ d ← ω ≠, z ◇ _ ← {p → d + ← ω ≠ p ← α[z, ← ω]} * ≡ ∼ ω ◇ d(Δ(-1+d)† ∼ 0 1 ⊢ φ z)}
  D2P ← {0 ≠ ω:Θ ◇ p → 2{p[ω] ← α[α ⊥ ω]} / ∼ ∘ □ ω → p ← ι ≠ ω}

```

This code is used in chunk 7.

4 Testing

We use the APLUnit testing framework to facilitate our testing of the Co-dfns compiler. The test harness is designed around a testing philosophy in which we ever only write black-box tests that work on the whole compiler using inputs that could be created or are expected to be creatable by end-users. That is, we do no “unit testing” of our source code, but only whole program testing.

The testing framework is provided by the `ut.apln` file, which is not part of this literate program and so is not included in this document. In order to make some of the testing more convenient, we define the function `TEST` to run the tests that exist in the `tests\` sub-directory. Each of these tests has a specific number which defines the test, and we refer to the tests by number when running them. Both of these testing functions assume that we are running inside of the `tests\` directory or one configured identically to it.

The `TEST` function takes either `'ALL'` as its input or a test number in the form of an integer. Given an integer, we call the test matching that number in the current working directory.

The `'ALL'` option causes `TEST` to run all of the tests that are defined in the current working directory. This command is a nicety, since we can technically do all of this by iterating the `TEST` function over the range of test numbers, but this would not create the aggregate statistics that we would like to see at the end of the testing report. By using `'ALL'` we get to see a complete summary of the results of testing all the code, rather than just the individual testing results on a per testing group/number basis.

```
16a <TEST 16a>≡
    TEST←{
      #.UT.(print_passed print_summary)←1
      'ALL'≡ω:#.UT.run './'
      path←'./t',(1 0⌞(4p10)⌞ω),'*_tests.dyalog'
      #.UT.run ⍵0⌞NINFO⌞1⌞path
    }
```

Root chunk (not used in this document).

Defines:

`TEST`, used in chunks 16b and 64a.

The `TEST` function is part of the utilities that exist outside of the `codfns` namespace, so we define a file for it.

```
16b <Tangle Commands 8>+≡
    echo "Tangling src/TEST.aplf..."
    notangle -R'[[TEST]]' codfns.nw > src/TEST.aplf
```

This code is used in chunk 83.

Defines:

`TEST.aplf`, never used.

Uses `codfns 7`, `src 88`, and `TEST 16a`.

5 Co-dfns Compiler

5.1 Parser

The first, and in many ways, the most complex element of the compiler is the parser. APL has a number of unique issues when it comes to adequately parsing the language, but the most important is handling the context-sensitive nature of parsing variables: depending on the type of a variable, the parse tree can look very different. To manage this, we make use of a linear, multi-pass style of parser in which the parsing process consists of numerous small passes over the input, each time refining the input into something more like the final result. The parser should take some input that matches the input requirements of the `Fix` function and produce a suitable output AST.

$$PS :: Source \rightarrow AST \times ExportTypes \times SymbolTable \times Source$$

We can think of the parser as starting with a forest of trees, each of which contains a single root node that represents a single character in from the input source, with all trees arranged in the source order. During each pass of the parser, we progressively combine these trees into more complex trees until we end up at the end with a single tree per parsed module. In other words, we take a fully flat forest of single-node trees and progressively increase the depth while reducing the number of root-nodes until we have our desired AST structure.

We divide the parsing roughly into two main phases, the tokenization phase and the parsing phase. Unlike most compilers, we don't have a strict division in these two phases, so, as they say, think of them more like guidelines than actual rules⁴.

```

17  ⟨Parser 17⟩≡
    PS←{
      ⟨Verify source input ω, set IN 14a⟩

      ⟨Parsing Constants 18a⟩
      ⟨Line and error reporting utilities 20b⟩

      ⟨Tokenize input 21⟩
      ⟨Parse token stream 22⟩

      ⟨Compute parser exports 66b⟩
      ⟨Adjust AST for output 18b⟩
    }

```

⁴<https://www.youtube.com/watch?v=WJVBvvS57j0>

This code is used in chunk 7.

Defines:

PS, used in chunks 13 and 88.

When parsing, it's very helpful to have names for line endings.

18a $\langle \text{Parsing Constants } 18a \rangle \equiv$

CR LF \leftarrow UCS 13 10

This code is used in chunk 17.

5.1.1 Output of the Parser

After we finish all of our parsing, we need to take the resulting AST and convert that into something that is suitable for output to the rest of the system. We do this in a few ways.

When we finish parsing, we expect the following fields:

Field	Description
d	Depth vector
t	Node type
k	Node sub-class or "kind"
n	Name/value field
pos	Starting index for source position
end	Exclusive index for source end position
xn	Names of top-level exported bindings
xt	Types of top-level exported bindings
sym	Symbol Table
IN	Canonical source code

On parser output, we want to convert the AST to an order that follows a depth-first, preorder traversal order, so that we can switch from using the parent vector to the depth vector. We use this output as our main output because it is space efficient for storage, and it works well as a canonical form to use. Because applications may want to only use the parser and not the rest of the compiler, we want to choose an output format that is suitable for external as well as internal use. This has some performance overheads, but it is probably worth it regardless, as reordering at this point to allow a depth vector enables some nice assumptions in the rest of the compiler. We use the P2D utility to reorder all of our AST columns. Note that things like the exported bindings and the symbol table are not strictly part of the AST structure, because they are of a different length and type than the other columns.

18b $\langle \text{Adjust AST for output } 18b \rangle \equiv$

d i \leftarrow P2D p \diamond d n t k pos end I \rightarrow \leftarrow c i

This definition is continued in chunks 19 and 20a.

This code is used in chunk 17.

There is an inefficiency in the AST representation at this point, where the `n` field contains character vectors. This inefficiency was necessary while building up the AST because we were not sure what symbols would be created before we parsed them, but at this point, we know the full set of symbols that we have in the AST. This means that we can convert the `n` field to a symbol table representation. In this case, we want the `n` field to pair with a `sym` list that contains all the unique symbols in the source. We want `old_n ≡ sym[|new_n]` to hold for this new `n` field. In other words, we want the new `n` field to contain negative integers whose magnitudes are valid indices into the `sym` symbol table. This means that there is only one character vector per unique symbol or numeric literal in the source code, which can greatly reduce memory usage. Moreover, it is much faster to compare symbols that are represented by numeric index rather than character vector. Most of the work we expect to be done on the `n` field, so that we never have to pull in `sym` unless we want to know the actual value of the symbol. This actually mimics the feature of symbols in other languages like Scheme, but it comes with an additional efficiency benefit in that we do not require the use of a full generalized pointer to represent a symbol if we have fewer symbols. This means that we are very likely only going to need a single byte or a couple of bytes per symbol to represent it in the `n` field.

The choice to make all of our symbols negative in value is somewhat strange, but we have a good reason for doing so. The `n` field is a single field that we use to contain general data for every node, and as such, it represents a sort of union type of all sorts of different data. In particular, we also want to be able to support using the `n` field to point to other nodes in the AST, which is a feature we rely heavily on in the compiler transformations. However, this feature would conflict with using the `n` field as an index into the `sym` table, rather than as an index into the AST. By making symbol pointers negative, we put them into a separate space than the positive AST node pointers, allowing us to store both pointers in the same field. This may seem like a little bit of a strange hack, but it actually makes reasoning about things a little easier, because we can tend to think of `n` as a name, even if that name is pointing to an AST or a symbol, and avoids needless space duplication or the need to remember to update multiple fields that are only relevant for some nodes.

We map the 0th index to be a null or empty symbol. We also want to reserve the first four symbol slots [1, 4] so that they will *always* refer to the same symbols, namely, ω , α , $\alpha\alpha$, and $\omega\omega$.

This gives us the following definitions for `sym` and `n`.

```
19 (Adjust AST for output 18b) +=
    sym ← v('')(, 'ω')(, 'α') 'αα' 'ωω', n
    n ← -sym ∷ n
```

This code is used in chunk 17.

Finally, we want to return our AST structure in a meaningful way. Logically, we have the AST proper, which consists of these fields:

```
d t k n pos end
```

The above fields are returned as an inverted table, where each column is a vector of the same length. We also want to return the variable environment, which gives the names of our top-level bindings and their types, also as an inverted table. Finally, we must return a canonical representation of the source code that is suitable as an indexing target for the `pos` and `end` fields, as well as the symbol table. Thus, we have a four element vector as the return value:

```
AST TopBindingTypes SymbolTable InputSource
```

Which gives us the following return value.

20a *⟨Adjust AST for output 18b⟩* \equiv
`(d t k n pos end)(xn xt)sym IN`

This code is used in chunk 17.
 Uses `xn` 66b and `xt` 66b.

5.1.2 Handling Parsing Errors

20b *⟨Line and error reporting utilities 20b⟩* \equiv

```

linestarts←(⌊1;2>∇IN∈CR LF)∇≠IN
mkdm←{α+2 ⋄ line←linestarts⌊ω ⋄ no←['',(⌘1+line),'] '
i←(∼IN[i]∈CR LF)∇i←beg+⌊linestarts[line+1]-beg←linestarts[line]
(□EM α)(no,IN[i])(' ^'[i∈ω],⌘' 'ρ⌘≠no)}
quotelines←{
lines←⌊linestarts⌊ω
nos←(1 0ρ⌘2×≠lines)∇['',(⌘1+lines),⌘1⌊'] '
beg←linestarts[lines] ⋄ end←linestarts[lines+1]
m←εω⌘i←beg+⌊end-beg
⌘1↓enos,(∼ωCR LF⌘⌘,(IN∘I⌘i),⌘' ⌘'∘I⌘m),CR}
SIGNAL←{α+2 ' ' ⋄ en msg←α ⋄ EN∘←en ⋄ DM∘←en mkdm ∇ω
dmx←('EN' en)('Category' 'Compiler')('Vendor' 'Co-dfns')
dmx,←c'Message'(msg,CR,quotelines ω)
□SIGNAL←dmx}

```

This code is used in chunk 17.

Defines:

```

linestarts, never used.
mkdm, never used.
quotelines, used in chunks 47c and 49a.
SIGNAL, used in chunks 14a, 24–27, 47–49, 52–54, 58c, 60a, 61c, 63–66, 68c, 78,
and 81c.
Uses dmx 43a.

```

5.1.3 Tokenizing the Input

```

21  <Tokenize input 21>≡
    A Group input into lines as a nested vector
    pos←(1≠IN)⊆~IN∈CR LF

    <Check and mask the strings 49a>
    <Unify whitespace and comments 48a>
    <Tokenize strings 49b>
    <Verify that all open characters are valid 47c>
    <Tokenize numbers 48b>
    <Tokenize variables 49c>
    <Tokenize primitives and atoms 51b>
    <Compute dfns regions and type, with } as a child 61c>
    <Check for out of context dfns formals 49d>
    <Compute trad-fns regions 63c>
    <Identify label colons vs. others 64d>
    <Tokenize keywords 65a>
    <Tokenize system variables 51d>

    A Delete all characters we no longer need from the tree
    d tm t pos end(⌈~)←c(t≠0)∨x∈'()[\{}:; '

    <Tokenize labels 64e>

```

This code is used in chunk 17.

5.1.4 Parsing Token Stream

22 *⟨Parse token stream 22⟩*≡

A Now that all compound data is tokenized, reify n field before tree-building

```
n←{1↓⊥''0',ω}@{t=N}{c''}@{t∈Z F}1 ⊞C@{t∈K S}IN∘I''pos+i''end-pos
```

⟨Check that all keywords are valid 65b⟩

⟨Check that namespaces are at the top level 65c⟩

⟨Verify that all structured statements appear within trad-fns 68c⟩

⟨Verify that system variables are defined 52a⟩

A Compute parent vector from d

```
p←D2P d
```

⟨Compute the nameclass of dfns 61d⟩

A We will often wrap a set of nodes as children under a Z node

```
gz←{
z←ω↑⊥-0≠ω ⊞ ks←-1↓ω
t[z]←Z ⊞ p[ks]←z ⊞ pos[z]←pos[ω] ⊞ end[z]←end[ωz,ks]
z
}
```

⟨Nest top-level root lines as Z nodes 65d⟩

⟨Wrap all dfns expression bodies as Z nodes 61e⟩

A Drop/eliminate any Z nodes that are empty or blank

```
_←p[i]{msk[α,ω]←~^≠IN[pos[ω]]∈WS}⊞i←⊥(t[p]=Z)∧p≠i≠p-msk←t≠Z
tm n t k pos end(⊥)←msk ⊞ p←(⊥~msk)(⊥-1+⊥)msk≠p
```

⟨Parse :Namespace syntax 66a⟩

⟨Parse guards to (G (Z ...) (Z ...)) 64a⟩

⟨Parse brackets and parentheses into -1 and Z nodes 58c⟩

⟨Convert ; groups within brackets into Z nodes 52c⟩

⟨Parse Binding nodes 54a⟩

⟨Mark system variables as P nodes with appropriate kinds 52b⟩

⟨Mark atoms, characters, and numbers as kind 1 50b⟩

⟨Mark APL primitives with appropriate kinds 51c⟩

⟨Anchor variables to earliest binding in the matching frame 61f⟩

⟨Convert M nodes to F0 nodes 68d⟩

⟨Convert α and ω to V nodes 49e⟩

⟨Convert αα and ωω to P2 nodes 49f⟩

⟨Infer the type of bindings, groups, and variables 54b⟩

⟨Strand arrays into atoms 50c⟩

⟨Parse dyadic operator bindings 54c⟩

⟨Rationalize F[X] syntax 53e⟩

(Group function and value expressions 58d)
(Parse function expressions 60a)
(Parse assignments 55c)
(Enclose $V[X; \dots]$ for expression parsing 53a)
(Parse trains 60c)
(Parse value expressions 59b)
(Rationalize $V[X; \dots]$ 53b)

```

A Sanity check
ERR←'INVARIANT ERROR: Z node with multiple children'
ERR assert(+/t[p]=Z)^p≠i≠p)=+/t=Z:

A Count parentheses in source information
ip←p[i←1(t[p]=Z)^n[p]∈c, '('] ♦ pos[i]←pos[ip] ♦ end[i]←end[ip]

A VERIFY Z/B NODE TYPES MATCH ACTUAL TYPE

A Eliminate Z nodes from the tree
zi←p I@{t[p[ω]]=Z}*≡ki←1msk←(t[p]=Z)^t≠Z
p←(zi@ki≠p)[p] ♦ t k n pos end(¬@zi)←t k n pos end I''cki
t k n pos endf''←msk←~mskv t=Z ♦ p←(1~msk)(t-1+1)mskf p

```

This code is used in chunk 17.

Uses assert 78.

5.2 Compiler Transformations

23 *(Compiler 23)≡*

```

TT←{
  ((d t k n ss se)exp sym src)←ω

  A Compute parent vector and reference scope
  r←I@{t[ω]≠F}*≡p→2{p[ω]←α[α1ω]}f t o c d→p←i≠d

  (Lift dfns to the top-level 62a)
  (Wrap expressions as binding or return statements 62b)
  (Lift guard tests 64b)
  (Count strand and indexing children 50d)
  (Lift and flatten expressions 59a)
  (Compute slots and frames 62d)
  (Record exported top-level bindings 66c)

  p t k n f s r d xi sym
}

```

This code is used in chunk 7.

Uses src 88 and xi 66c.

5.3 Code Generator

24a *⟨Map generators over the linearized AST; return 24a⟩*≡
`d i←P2D p ♦ ast←(Q†d p t k n(ι≠p)fr sl fd)[i;] ♦ ks←{ω<[0]~(▷ω)=ω[;0]}`
`NOTFOUND←{(' [GC] UNSUPPORTED NODE TYPE ',NΔ[▷ω],⌘φω)⊠SIGNAL 16}`
`dis←{0=2▷h←,1†ω:' ' ♦ (≠gck)=i←gckι<h[2 3]:NOTFOUND h[2 3] ♦ h(±i▷gcv)ks 1†ω}`
`ε,◦(⊠UCS 13 10)``pref,▷,†(,†Zp``t=F),(,†Zx``xi),(c<''),dis``ks ast`

This code is used in chunk 25b.
 Uses SIGNAL 20b and x i 66c.

24b *⟨Symbol ↔ Name mapping 24b⟩*≡
`syms←0p<' ' ♦ nams←0p<' '`
 This definition is continued in chunks 53c, 56a, 62e, and 68–77.
 This code is used in chunk 25b.

24c *⟨Node ↔ Generator mapping 24c⟩*≡
`gck←0p<0 0 ♦ gcv←0p<' '`
 This definition is continued in chunks 48–50, 53d, 55a, 56b, 59c, 60b, 62, 64c,
 and 66d.
 This code is used in chunk 25b.

24d *⟨Prefix code for all generated files 24d⟩*≡
`pref ←<'#include "codfns.h"'`
`pref,←<' '`
`pref,←<'EXPORT int'`
`pref,←<'DyalogGetInterpreterFunctions(void *p)'`
`pref,←<{'`
`pref,←<' return set_dwafns(p);'`
`pref,←<'}'`
`pref,←<' '`
 This code is used in chunk 25b.
 Uses codfns 7, codfns.h 33, and set_dwafns 46a.

24e *⟨Node-specific code generators 24e⟩*≡
`Zp←{`
`n←'fn',⌘ω`
⟨Declare top-level function bindings 61a⟩
`'UNKNOWN FUNCTION TYPE'⊠SIGNAL 16`
`}`
 This definition is continued in chunks 25a, 50a, 51a, 55b, 59d, 62, 63, and 67.
 This code is used in chunk 25b.
 Uses SIGNAL 20b.

25a $\langle \text{Node-specific code generators 24e} \rangle + \equiv$

```

  Zx ← {
    n ← sym ▷ n[ω]  ◊  rid ← rrf[ω]
    k[ω] = 0 : c ' '
     $\langle \text{Declare top-level array structures 50f} \rangle$ 
     $\langle \text{Declare top-level closures 61b} \rangle$ 
    ⚡ ' ' UNKNOWN EXPORT TYPE ' ' □ SIGNAL 16 '
  }

```

This code is used in chunk 25b.
Uses EXPORT 34b and SIGNAL 20b.

25b $\langle \text{Code Generator 25b} \rangle \equiv$

```

  GC ← {
    p t k n fr sl rf fd xi sym ← ω

     $\langle \text{Symbol} \leftrightarrow \text{Name mapping 24b} \rangle$ 
     $\langle \text{Node} \leftrightarrow \text{Generator mapping 24c} \rangle$ 

     $\langle \text{Prefix code for all generated files 24d} \rangle$ 
     $\langle \text{Node-specific code generators 24e} \rangle$ 

     $\langle \text{Map generators over the linearized AST; return 24a} \rangle$ 
  }

```

This code is used in chunk 7.
Uses xi 66c.

5.4 Backend C Compiler Interface

26 *(Interface to the backend C compiler 26)≡*

```

CC←{
  vsbat←VSΔPATH, '\VC\Auxiliary\Build\vcvarsall.bat'
  soext←{opsys'.dll' '.so' '.dylib'}
  libdir←opsys '' '/lib64' '' '/lib' ''
  ccf←{' -o ',ω, '.',α, '' ''',ω, '.c' ' -laf',AFΔLIB, ' > ',ω, '.log 2>&1'}
  cci←{'-I'',AFΔPREFIX, '/include' ' -L'',AFΔPREFIX, libdir}
  cco←'-std=c99 -Ofast -g -Wall -fPIC -shared '
  cco,←'-Wno-parentheses -Wno-misleading-indentation '
  ucc←{ωω(□SH αα, ' ',cco,cci,ccf)ω}
  gcc←'gcc'ucc'so'
  clang←'clang'ucc'dylib'
  vsco←{z←'/W3 /wd4102 /wd4275 /O2 /Zc:inline /Zi /FS /Fd"',ω, '.pdb' '
  z,←'/WX /MD /EHsc /nologo '
  z, '/I"%AF_PATH%\include" /D "NOMINMAX" /D "AF_DEBUG" '}
  vslo←{z←'/link /DLL /OPT:REF /INCREMENTAL:NO /SUBSYSTEM:WINDOWS '
  z,←'/LIBPATH:"%AF_PATH%\lib" /OPT:ICF /ERRORREPORT:PROMPT /TLBID:1 '
  z, '/DYNAMICBASE "af', AFΔLIB, '.lib" "codfns.lib" '}
  vsc0←{~□NEXISTS vsbat:'VISUAL C?'□SIGNAL 99 ♦ '""',vsbat, '" amd64'}
  vsc1←{' && cd "',(□CMD'echo %CD%'),' && cl ',(vsco ω), ' "',ω, '.c' ' '}
  vsc2←{(vslo ω), '/OUT:"',ω, '.dll' > '"',ω, '.log'""'}
  vsc←{□CMD ('%comspec% /C ',vsc0,vsc1,vsc2)ω}
  _←(±opsys'vsc' 'gcc' 'clang')α→ω put α, '.c'→1 □NDELETE f←α,soextθ
  □←,→□NGET(α, '.log')1
  □NEXISTS f:f ♦ 'COMPILE ERROR' □SIGNAL 22}

```

This code is used in chunk 7.

Uses AFΔLIB 11, AFΔPREFIX 11, codfns 7, opsys 82b, put 81c, SIGNAL 20b, vsbat 89a, vsc 89a, and VSΔPATH 12.

5.5 Linking with Dyalog

27 *(Linking with Dyalog 27)*≡

```

NS←{
  MKA←{mka←ω} ◇ EXA←{exa θ ω}
  Display←{α←'Co-dfns' ◇ W←w_new←α ◇ 777::w_del W
  w_del W←W αα{w_close α:±'□SIGNAL 777' ◇ α αα ω}*ωω←ω}
  LoadImage←{α←1 ◇ ~□NEXISTS ω:□SIGNAL 22 ◇ loading θ ω α}
  SaveImage←{α←'image.png' ◇ saveimg ω α}
  Image←{~2 3v.=≠pω:□SIGNAL 4 ◇ (3≠pω)^3=≠pω:□SIGNAL 5 ◇ ω←w_img ω α}
  Plot←{2≠pω:□SIGNAL 4 ◇ ~2 3v.=1pω:□SIGNAL 5 ◇ ω←w_plot (⊞ω) α}
  Histogram←{ω←w_hist ω,α}
  RtmΔInit←{
    _←'w_new'□NA'P ' ,ω,'|w_new <C[]'
    _←'w_close'□NA'I ' ,ω,'|w_close P'
    _←'w_del'□NA ω,'|w_del P'
    _←'w_img'□NA ω,'|w_img <PP P'
    _←'w_plot'□NA ω,'|w_plot <PP P'
    _←'w_hist'□NA ω,'|w_hist <PP F8 F8 P'
    _←'loading'□NA ω,'|loading >PP <C[] I'
    _←'saveimg'□NA ω,'|saveimg <PP <C[]'
    _←'exa'□NA ω,'|exarray >PP P'
    _←'mka'□NA'P ' ,ω,'|mkarray <PP'
    _←'FREA'□NA ω,'|frea P'
    _←'Sync'□NA ω,'|cd_sync'
    0 0 ρ θ}
  mkna←{α,'|',('Δ'□R'___'←ω),'_cdf P P P'}
  mkf←{
    fn←α,'|',('Δ'□R'___'←ω),'_dwa '
    z←c'Z←{A}' ,ω,' W'
    z,←c':If 0=□NC'Δ.' ,ω,'_mon'''
    z,←c'      ' ,ω,'_mon'Δ.□NA''' ,fn,'>PP P <PP'''
    z,←c'      ' ,ω,'_dya'Δ.□NA''' ,fn,'>PP <PP <PP'''
    z,←c':EndIf'
    z,←c':If 0=□NC'A'''
    z,←c'      Z←Δ.' ,ω,'_mon 0 0 W'
    z,←c':Else'
    z,←c'      Z←Δ.' ,ω,'_dya 0 A W'
    z,←c':EndIf'
    z
  }
  ns←#.□NSθ ◇ _←'ΔΔ'ns.□NS''cθ ◇ Δ Δ←ns.(Δ Δ)
  Δ.names←(0p<''),(2=1>α)≠0>α
  fns←'RtmΔInit' 'MKA' 'EXA' 'Display'
  fns,←'LoadImage' 'SaveImage' 'Image' 'Plot' 'Histogram'
  fns,←'soext' 'opsys' 'mkna'

```

```

_←Δ.⊞FX⊞CR``fns
Δ.(decls←ω⊞mkna``names)
_←ns.⊞FX``(c''),ω⊞mkf``Δ.names
_←'Z←Init'
_,←c'Z←RtmΔInit'',ω,``''
_,←c'→0/˜0=≠names'
_,←c'names ##.Δ.⊞NA``decls'
_←Δ.⊞FX _
ns
}

```

This code is used in chunk 7.
 Uses PP 81a and SIGNAL 20b.

```

29  (DWA Function Export 29)≡
    z,←c'EXPORT int'
    z,←cn,'_dwa(struct localp *zp, struct localp *lp, struct localp *rp)'
    z,←c'{'
    z,←c'    struct array *z, *l, *r;'
    z,←c'    int err;'
    z,←c'
    z,←c'    l = NULL;'
    z,←c'    r = NULL;'
    z,←c'
    z,←c'    fn',rid,'(NULL, NULL, NULL, NULL);'
    z,←c'
    z,←c'    err = 0;'
    z,←c'
    z,←c'    if (lp)'
    z,←c'        err = dwa2array(&l, lp->pocket);'
    z,←c'
    z,←c'    if (err)'
    z,←c'        dwa_error(err);;'
    z,←c'
    z,←c'    if (rp)'
    z,←c'        dwa2array(&r, rp->pocket);'
    z,←c'
    z,←c'    if (err) {'
    z,←c'        release_array(l);'
    z,←c'        dwa_error(err);'
    z,←c'    }'
    z,←c'
    z,←c'    err = (' ,n,'->fn)(&z, l, r, ' ,n,'->fv);'
    z,←c'
    z,←c'    release_array(l);'
    z,←c'    release_array(r);'
    z,←c'
    z,←c'    if (err)'
    z,←c'        dwa_error(err);'
    z,←c'
    z,←c'    err = array2dwa(NULL, z, zp);'
    z,←c'    release_array(z);'
    z,←c'
    z,←c'    if (err)'
    z,←c'        dwa_error(err);'
    z,←c'
    z,←c'    return 0;'
    z,←c'}'
    z,←c'

```

This code is used in chunk 61b.

July 10, 2022

codfns.nw 30

Uses dwa_error 44a.

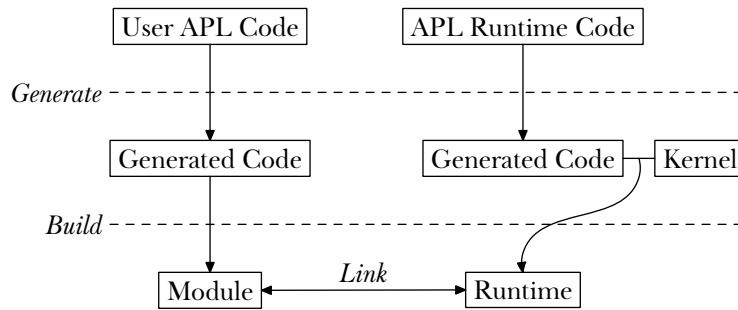


Figure 1: Process of Building and Linking the Runtime

5.6 Runtime

The runtime component of Co-dfns handles the code necessary for the output of the Code Generator to run. This includes support for all the supported language features as well as the runtime code for the built-in APL primitives and system functionality. The design of the runtime is meant to allow for as much of the runtime as possible to be implemented in APL. We also want to make it as easy as possible to target new languages for output from the compiler.

Conceptually, the code generator produces a code module that links against an already built runtime module that provides all the language support. Each module has some “backend target” language. In order to make retargeting the compiler as simple as possible and to implement most of the runtime as APL, we split the runtime code into an APL namespace, containing all the APL code that is applicable to all backends and that can be implemented in APL, and a backend kernel that contains all the backend language-specific code that we must use. We can split the compiler into a frontend *generate* and a backend *build* step. The generate phase takes the input APL source and generates code in the backend target language that depends on a runtime implementation. The build phase takes that code and uses the backend toolchain to link, compile, and otherwise assemble the code into an appropriate redistributable “binary”. The C backend, for instance, takes APL and turns it into C code where a C compiler then builds and links it against a runtime, finally producing a DLL.

To build the runtime, the same basic approach is used. We use the compiler to generate a backend file from the APL runtime code. However, since no runtime exists for the runtime itself, we do *not* continue in the typical manner and build with the standard backend pipeline, which assumes the existence of a runtime. Instead, we merge the generated code with the kernel for that specific backend and build as its own standalone object.

This workflow is illustrated by Figure 1 showing how all of the

pieces of the runtime interact with user code.

This architecture has some interesting advantages. First, most of the process for building the runtime is just like building any other piece of APL code. Second, only a small kernel and code generator need to be implemented for a new backend, with most of the work remaining in the APL runtime code. Third, the runtime may be implemented using a different backend language than that used for compiling the user code. All that is required is that the backend for the user code knows how to link to and access the code in the runtime object. This permits, for instance, a Scheme or Javascript backend to depend on a runtime implemented in C, thus enabling greater performance while hiding any integration hassles from the interface exposed by the user module. In theory, any combination of suitable backend languages may be used.

We put all the runtime primitives into a single Co-*dfns* namespace called `prim.apln`.

```
32a  <prim.apln 32a>≡
      :Namespace prim

      <APL Primitives 76f>
      <System Primitives (never defined)>

      :EndNamespace

Root chunk (not used in this document).
Defines:
  prim, used in chunks 32b and 88.
  prim.apln, used in chunk 32b.

32b  <Tangle Commands 8>+≡
      echo "Tangling rtm/prim.apln..."
      notangle -R'prim.apln' codfns.nw > rtm/prim.apln
```

This code is used in chunk 83.
Uses `codfns 7`, `prim 32a`, and `prim.apln 32a`.

Each primitive has its own unique considerations, so we leave the definition of these primitives to section 7.

For each backend we must have a unique kernel and code generator. Most of that content will be defined on a per-language feature basis below. The rest of this section focuses on the more generic and fundamental elements of the kernels, such as general organization, interface, and memory management.

5.6.1 GPU C Runtime Kernel

The main concern of a C runtime is managing memory and adequately handling access to the DWA system. Dyalog's DWA system permits us direct access to the underlying interpreter array format and memory manager. We could use this format directly but this will not work for GPU compute because the DWA interface connects array elements and header information in a way that makes GPU allocating them quite difficult, especially if we only want the elements on the GPU.

DWA has a specific array format, but we will delay specifying utility code for array handling until section 6.6. In this section, we handle the following issues:

- DWA Initialization
- Header Structure
- Memory Management
- Datatype Management
- Error Reporting

We deal with the top-level error signalling behavior in this section, but for error signalling within functions, as well as arrays, module initialization function calls, and so forth, see the appropriate subsection of Language Features (section 6).

The first order of business is the main structure of the C runtime files and API. We could attempt to put all our runtime code into a single `kernel.c` file, but the result would require us to maintain includes in a way that prevents us from easily linking the include statements to each language feature implementation without encouraging needless duplicate includes. Instead, we assume that each language feature will be given its own C file and then we can manage includes independently. We will make use of a single `codfns.h` file that contains all the public entry points into the runtime.

33 `<codfns.h 33>=`
`#pragma once`

```

<C runtime includes (never defined)>
<C runtime macros 34b>
<C runtime enumerations 36b>
<C runtime structures 36a>
<C runtime declarations 38a>

```

Root chunk (not used in this document).

Defines:

`codfns.h`, used in chunks 24d, 34a, 41a, 58a, 68a, and 89b.

```

34a <Tangle Commands 8>+≡
    echo "Tangling rtm/codfns.h..."
    notangle -R'codfns.h' codfns.nw > rtm/codfns.h

```

This code is used in chunk 83.

Uses `codfns` 7 and `codfns.h` 33.

Since we want to use this single header for the runtime code *and* the generated code that will import the runtime, an interesting situation arises regarding exports. Both generated and runtime code must export functions from their respective DLLs, but in the case of the runtime, these exported functions are also the functions that we must import into our generated code, we must annotate the edeclaration of such functions differently if we are importing than when we are exporting. Thus, when we are building the runtime, we want to export all our bindings, but when we are accessing the runtime from generated code we want to import those same bindings while exporting functions that we generate.

To handle this, we rely on three preprocessor definitions. When we are building the runtime, we will define `EXPORTING`, but we expect this to be undefined when building generated code. Then we have an `EXPORT` definition that always maps to the platform specific export decorator, while `DECLSPEC` will be the import spec or export spec depending on `EXPORTING`.

It used to be the case that each platform handled DLL importing and exporting differently, but modern compilers all handle the `__declspec` syntax, so we will use that for all platforms.

```

34b <C runtime macros 34b>≡
    #define EXPORT __declspec(dllexport)
    #ifdef EXPORTING
    #define DECLSPEC EXPORT
    #else
    #define DECLSPEC __declspec(dllimport)
    #endif

```

This code is used in chunk 33.

Defines:

`DECLSPEC`, used in chunks 37–40, 43, 44, 46a, and 57.

`EXPORT`, used in chunks 25a and 68a.

`EXPORTING`, used in chunk 89a.

Our next major concern is handling memory and multiple data types. Since the compiler assumes a stack machine model, we have a unified stack that will contain many different objects, such as functions and arrays, so we must have a way of handling the objects in a somewhat generic way.

While some generality is desirable, I must curtail my Scheme-esque impulse towards unnecessary dynamic generality. This is a runtime, after all, and experience shows that extra dynamic annotation can seriously impede scalability of the system and introduce unfortunate performance gotchas. Rather than chase this form of programmability, I am taking a page from Knuth's book and aiming for "re-editable" code that can be easily, but statically, extended. The goal is to avoid excess runtime allocation and indirection while at the same time making it easy to add and manage datatypes.

Any such memory or type management system must address the following questions:

- How do I make an object?
- How do I free an object?
- When do I free an object?
- How do I keep an object alive?
- How do I make new data types?

In APL, most values have a stack lifetime, which would encourage us to make use of a stack semantics in our runtime. However, for more involved APL, this assumption does not hold true. Instead, to manage our objects, we choose to make use of reference counting.⁵ This maintains most of the predictability and low-overhead of a stack semantics but gives us the additional power to allow object lifetimes to extend beyond the lifetime of their definition context.

We do not have a requirement in our system for generic object creation (indeed, such a requirement is quite rare), but we do need to generically retain a reference to an object and to release an object. We want to enable this without too much indirection. To implement this, we simply require that all our datatypes be structures that share the following common fields. We call these types cells as a convenient term.

35 *⟨Common cell fields 35⟩*≡
 enum cell_type ctyp;
 unsigned int refc;

This code is used in chunks 36a and 56d.
 Defines:

⁵https://en.wikipedia.org/wiki/Reference_counting

`ctyp`, used in chunks 37, 39a, and 57a.
`refc`, used in chunks 37, 38b, 40a, and 57a.
 Uses `cell_type` 36b.

These fields help us to answer the two most important questions we must answer for any cell: what type of cell is it; and, is it currently referenced? By requiring all data structs to have these fields in common, we can cast them about and be basolutely sure that things will continue to work. We define a “void” cell type `struct cell_void` to be our minimal cell type.

36a *⟨C runtime structures 36a⟩*≡
 `struct cell_void {`
 ⟨Common cell fields 35⟩
 `};`

This definition is continued in chunk 56d.

This code is used in chunk 33.

Defines:

`cell_void`, used in chunks 37–40.

The `enum cell_type` keeps track of all known cell types.

36b *⟨C runtime enumerations 36b⟩*≡
 `enum cell_type {`
 ⟨Cell type names 36c⟩
 `};`

This code is used in chunk 33.

Defines:

`cell_type`, used in chunk 35.

We set the first 0th cell type to our void cell.

36c *⟨Cell type names 36c⟩*≡
 `CELL_VOID`

This definition is continued in chunk 56c.

This code is used in chunk 36b.

Defines:

`CELL_VOID`, used in chunks 37 and 39c.

We do not make or define any generic way to create cells; you must make a constructor function suitable to the needs of the data type. At the moment, it is the responsibility of such makers to ensure that the common fields are appropriately initialized. A maker should return a 0 on success and a non-zero error on failure. It should also take a `struct cell_TYPE **` as the first argument to store the allocated cell in. We expect the slot passed to a creator will be a possibly previously utilized slot on a stack or something along these lines. This means that it is the caller's responsibility to ensure that this slot has already been released. Failure to do this would potentially lead to a memory leak. However, attempting to handle this within the cell maker function results in an API that is much too fragile and needlessly complex. We expect to generally follow the stylistic guideline that a function should allocate and own its own data and then release that data in the same function.

The basic cell maker for the `void` cell type looks like this:

```
37  <Cell definitions 37>≡
    DECLSPEC int
    mk_void(struct cell_void **cell)
    {
        struct cell_void *ptr;

        ptr = malloc(sizeof(struct cell_void));

        if (ptr == NULL)
            return 1;

        ptr->ctyp = CELL_VOID;
        ptr->refc = 1;
        *cell = ptr;

        return 0;
    }
```

This definition is continued in chunks 38–40.

This code is used in chunk 41a.

Defines:

`mk_void`, used in chunk 38a.

Uses `CELL_VOID` 36c, `cell_void` 36a, `ctyp` 35, `DECLSPEC` 34b, and `refc` 35.

A few points of style here. The error codes should try to follow the standard APL codes. Additionally, the target slot should not be mutated until we are sure that all is well and that the object is well-formed.

38a *⟨C runtime declarations 38a⟩*≡
`DECLSPEC int mk_void(struct cell_void **);`

This definition is continued in chunks 38–40, 43, 44, and 57b.

This code is used in chunk 33.

Uses `cell_void` 36a, `DECLSPEC` 34b, and `mk_void` 37.

While we must define unique constructors for the various types, when releasing or freeing a cell of some kind, we *do* want to be able to generically free a cell. However, this must be done with a minimum of runtime overhead. First, we distinguish the terms “release” and “free”. If an object is freed, that object’s memory is fully returned to the memory manager, whereas releasing is about reducing the number of references to that object. When a cell has no references to it, then it is freed.

Each cell type will require its own unique release function that manages cleanly destroying the cell. The release function for the `void` cell type looks like this:

38b *⟨Cell definitions 37⟩*+≡
`DECLSPEC void`
`release_void(struct cell_void *cell)`
`{`
`if (cell == NULL)`
`return;`
`if (--cell->refc)`
`return;`
`free(cell);`
`}`

This code is used in chunk 41a.

Defines:

`release_void`, used in chunks 38c and 39c.

Uses `cell_void` 36a, `DECLSPEC` 34b, and `refc` 35.

38c *⟨C runtime declarations 38a⟩*+≡
`DECLSPEC void release_void(struct cell_void *);`

This code is used in chunk 33.

Uses `cell_void` 36a, `DECLSPEC` 34b, and `release_void` 38b.

To support generic cell release, we define a `release_cell` function.

```
39a  <Cell definitions 37>+≡
      DECLSPEC void
      release_cell(void *cell)
      {
        if (cell == NULL)
          return;

        switch (((struct cell_void *)cell)->ctyp) {
          <Cell release cases 39c>
          default:
            dwa_error(99);
          }
        }
      }
```

This code is used in chunk 41a.

Defines:

`release_cell`, used in chunks 39b and 57a.

Uses `cell_void` 36a, `ctyp` 35, `DECLSPEC` 34b, and `dwa_error` 44a.

```
39b  <C runtime declarations 38a>+≡
      DECLSPEC void release_cell(void *);
```

This code is used in chunk 33.

Uses `DECLSPEC` 34b and `release_cell` 39a.

For each cell type, we must plug the type-specific release function into this `release_cell` switch to enable generic releasing for that type. For the `void` type, this looks as follows:

```
39c  <Cell release cases 39c>≡
      case CELL_VOID:
        release_void(cell);
        break;
```

This definition is continued in chunk 57c.

This code is used in chunk 39a.

Uses `CELL_VOID` 36c and `release_void` 38b.

The above mostly suffices for dealing with cells. However, we also want to conveniently bump the reference count of a cell seamlessly without explicitly setting `refc`. We often encounter the case where we are assigning a cell to a new slot, thus requiring a reference count increment. The following function `retain_cell` lets us do this in a single statment by writing:

```
slot2 = retain_cell(slot1);
```

```
40a  <Cell definitions 37>+≡
      DECLSPEC void *
      retain_cell(void *cell)
      {
        if (cell != NULL)
          ((struct cell_void *)cell)->refc++;

        return cell;
      }
```

This code is used in chunk 41a.

Defines:

`retain_cell`, used in chunks 40b, 50a, 51a, and 55b.

Uses `cell_void` 36a, `DECLSPEC` 34b, and `refc` 35.

```
40b  <C runtime declarations 38a>+≡
      DECLSPEC void *retain_cell(void *);
```

This code is used in chunk 33.

Uses `DECLSPEC` 34b and `retain_cell` 40a.

Fortunately, this retention function requires no extra code as we extend the system with more data types. This gives us the following steps if we want to add a new data type to the runtime:

1. Add the cell type to *Cell type names 36c* as `, CELL_TYPE`.
2. Define the structure in *C runtime structures 36a*, making sure that *Common cell fields 35* are the first fields.
3. Define an `int mk_type(struct cell_type **, ...)` function and declare it in *C runtime declarations 38a*.
4. Define a `void release_type(struct cell_type *)` function and declare it in *C runtime declarations 38a*.
5. Add a case to *Cell release cases 39c* on `CELL_TYPE` that calls `release_type` on `cell`.

The cell handling we put into a file on its own.

41a *cell.c 41a*≡
`#include <stdlib.h>`

`#include "codfns.h"`

Cell definitions 37

Root chunk (not used in this document).

Defines:

`cell.c`, used in chunk 41b.

Uses `codfns 7` and `codfns.h 33`.

41b *Tangle Commands 8*+≡
`echo "Tangling rtm/cell.c..."`
`notangle -R'cell.c' codfns.nw > rtm/cell.c`

This code is used in chunk 83.

Uses `cell.c 41a` and `codfns 7`.

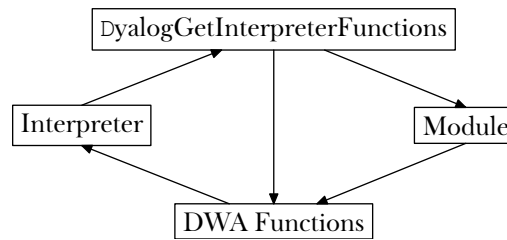


Figure 2: DWA module initialization

Finally, we must handle the DWA connection between a Co-dfns compiled module and the interpreter. One constraint on this design is the need to make a Co-dfns module work with or without a DWA-driven interpreter. If we are interfacing solely with a foreign, C-based system, we still must function somehow.

DWA modules export `DyalogGetInterpreterFns` as a function to link the interpreter and the module.

The function receives a structure from the interpreter populated with function pointers that enable access to various interpreter features. A small design point comes into play here because we do not want to unnecessarily expose our underlying model to the user of the compiled module. In particular, if an user is not a Dyalog interpreter, they should not need to know about the DWA system in order to function. For example, they should not need to know or use `DyalogGetInterpreterFunctions` or the underlying functions. Thus, we must have a way to achieve similar functionality from different systems.

Our approach to this is to provide more generic and explicit function for setting things we want from any system and then to layer DWA initialization on top of that.

Fundamentally, the main thing that we care about for all systems is having some means of making a non-local escaping error report. This main error reporting is meant to mimic the extended signalling functionality of the interpreter documented in the `DMX` object. The DWA equivalent of this structure is given by `struct dwa_dmx`.

42 *(DWA structures and enumerations 42)*≡

```

struct dwa_dmx {
    unsigned int flags;
    unsigned int en;
    unsigned int enx;
    const wchar_t *vendor;
    const wchar_t *message;
    const wchar_t *category;
};

```

This definition is continued in chunk 45.

This code is used in chunk 47a.

Defines:

`dwa_dmx`, used in chunks 43a and 44c.

In our APL model at the moment, there is only one main and universal `DMX` object at a time, so we define a single `dmx` binding to contain the current data.

43a *⟨DWA definitions 43a⟩*≡

```
struct dwa_dmx dmx;
```

This definition is continued in chunks 43, 44, and 46a.

This code is used in chunk 47a.

Defines:

`dmx`, used in chunks 20b, 43b, and 44a.

Uses `dwa_dmx` 42.

The reality of many FFI systems is that they do not do a good job of supporting C structs in the form of such global variables, so we must make sure that there is a meaningful way to access the system using nothing but function calls.

In the case of errors we have an interesting situation. In C, handling a long chain of errors demands that we are meticulous about how we handle the interaction of the call stack and any kind of early exit. In our case, this means that any time we finally call the non-local error function that we expect to never return, we may be quite far removed from the original site of the error. Thus, passing any complex data back up a call stack could be quite complex. Instead, we populate most of `dmx` that we care about using setter functions and then only have a very little to worry about passing up a call stack, namely, the error number itself.

we define a setter function `set_dmx_message` to handle setting `dmx.message`.

43b *⟨DWA definitions 43a⟩*+≡

```
DECLSPEC void
set_dmx_message(wchar_t *msg)
{
    dmx.message = msg;
}
```

This code is used in chunk 47a.

Defines:

`set_dmx_message`, used in chunk 43c.

Uses `DECLSPEC` 34b and `dmx` 43a.

43c *⟨C runtime declarations 38a⟩*+≡

```
DECLSPEC void set_dmx_message(wchar_t *);
```

This code is used in chunk 33.

Uses `DECLSPEC` 34b and `set_dmx_message` 43b.

Our main non-returning function `dwa_error` handles some of the parts of `dmx` that we do not currently change, and then calls the internally initialized error function provided by whatever our interfacing system is.

44a *⟨DWA definitions 43a⟩+≡*
`DECLSPEC void`
`dwa_error(unsigned int n)`
`{`
`dmx.flags = 3;`
`dmx.en = n;`
`dmx.enx = n;`
`dmx.vendor = L"Co-dfns";`
`dmx.category = NULL;`

`dwa_error_ptr(&dmx);`
`}`

This code is used in chunk 47a.

Defines:

`dwa_error`, used in chunks 29, 39a, and 44b.

Uses `DECLSPEC` 34b, `dmx` 43a, and `dwa_error_ptr` 44c.

44b *⟨C runtime declarations 38a⟩+≡*
`DECLSPEC void dwa_error(unsigned int);`

This code is used in chunk 33.

Uses `DECLSPEC` 34b and `dwa_error` 44a.

The above requires the calling interface set `dwa_error_ptr`, which we handle with `set_codfns_error`.

44c *⟨DWA definitions 43a⟩+≡*
`void (*dwa_error_ptr)(struct dwa_dmx *);`

`DECLSPEC void`
`set_codfns_error(void *fn)`
`{`
`dwa_error_ptr = fn;`
`}`

This code is used in chunk 47a.

Defines:

`dwa_error_ptr`, used in chunk 44a.

`set_codfns_error`, used in chunks 44d and 46b.

Uses `DECLSPEC` 34b and `dwa_dmx` 42.

44d *⟨C runtime declarations 38a⟩+≡*
`DECLSPEC void set_codfns_error(void *);`

This code is used in chunk 33.

Uses `DECLSPEC` 34b and `set_codfns_error` 44c.

To link this interface into the DWA functionality, we must extract the appropriate function pointers out of the structure passed to `DyalogGetInterpreterFunctions`. We assume that the code generator will create a suitable definition for `DyalogGetInterpreterFunctions` that calls the following `set_dwafns`, such as:

```
EXPORT int
DyalogGetInterpreterFunctions(void *fns)
{
    return set_dwafns(fns);
}
```

This established a link in each compiled module to the runtime DWA handling and allows us to keep the DWA logic inside the runtime. The DWA structure is relatively involved in its full expression, but we do not need the full power, so we can simplify our setup. We also want to talk about the structure more generically here without too much detail that may be more properly handled in the correct language feature section. At its heart, the structure is a set of functions, which we store as an array of `void *` pointers.

45 *(DWA structures and enumerations 42)+≡*

```
struct dwa_wsfns {
    long long size;
    void *fns[18];
};

struct dwa_fns {
    long long size;
    struct dwa_wsfns *ws;
};
```

This code is used in chunk 47a.

Defines:

```
dwa_fns, used in chunk 46a.
dwa_wsfns, never used.
```

It is the job of the `set_dwafns` function to set the appropriate Codfns interface functions and follow the initialization expectations of the DWA system. On successful initialization, the function should return 0, but we must check compatibility by examining the given structure `size`, return 16 if something is not right.

```
46a  <DWA definitions 43a>+≡
    DECLSPEC int
    set_dwafns(void *p)
    {
        struct dwa_fns *dwa;

        if (p == NULL)
            return 0;

        dwa = p;

        if (dwa->size < (long long)sizeof(struct dwa_fns))
            return 16;

        <Set DWA interface functions 46b>

        return 0;
    }
```

This code is used in chunk 47a.

Defines:

`set_dwafns`, used in chunk 24d.

Uses `DECLSPEC` 34b and `dwa_fns` 45.

Assuming that the DWA structure seems valid, we want to extract these functions into the appropriate names that we have created for them. An alternative would be to retain the structure and make indirect calls into that structure, but this is a little more awkward and would involve both more storage and more memory indirects for no more clarity and only more entanglement of the code. Instead, setting the correct names at the time of a `set_dwafns` call leads to a much cleaner dependency tree. At this point, only the `dwa_error` function has been designed and defined.

```
46b  <Set DWA interface functions 46b>≡
    set_codfns_error(dwa->ws->fns[17]);
```

This code is used in chunk 46a.

Uses `set_codfns_error` 44c.

47a $\langle \textit{dwa.c} \text{ 47a} \rangle \equiv$
 $\langle \textit{DWA includes}$ (never defined) \rangle
 $\langle \textit{DWA macros}$ (never defined) \rangle
 $\langle \textit{DWA structures and enumerations} \text{ 42} \rangle$
 $\langle \textit{DWA definitions} \text{ 43a} \rangle$
 Root chunk (not used in this document).
 Defines:
 `dwa.c`, used in chunk 47b.

47b $\langle \textit{Tangle Commands} \text{ 8} \rangle + \equiv$
 `echo "Tangling rtm/dwa.c..."`
 `notangle -R'dwa.c' codfns.nw > rtm/dwa.c`
 This code is used in chunk 83.
 Uses `codfns` 7 and `dwa.c` 47a.

6 Language Features

6.1 Valid source input character set

```

47c    <Verify that all open characters are valid 47c>≡
      alp←'ABCDEFGHIJKLMNOPQRSTUVWXYZ_abcdefghijklmnopqrstuvwxyz'
      alp,←'ÀÁÂÃÄÅÆÇÈÉÊËÌÍÎÏÐÑÒÓÔÕÖÙÚÛÜÝÞàáâãäåæçèéêëìíîïðñóôõöøùúûþ'
      alp,←'ΔΔABCDEFGHIJKLMNOPQRSTUVWXYZ'
      num←□D
      synb←'~[ ]{}()''':αω◇;'
      syna←'θ□□#'
```

prmfms←'+-×÷|/[*◎! ? ~^v~<≤=>≥≠≢ρ, ; φΘϙ↑↓↕≤▷∈∩υι⊥□△Ψ±∓⊥⊥→⇐⇒↔'

prmdo←'.°.*☒◻◻@' ◇ prmmo←'""~&I☒' ◇ prmfo←'/\¥'

prms←prmfms,prmdo,prmmo,prmfo

x←' '@{t≠0}IN[pos] A The spaces produce nice invariants

v/msk←~xεalp,num,syna,synb,prms,WS:{

EM←'SYNTAX ERROR: INVALID CHARACTER(S) IN SOURCE',CR

EM,←quotelines ⊆msk

EM □SIGNAL 2

}θ

This code is used in chunk 21.

Uses quotelines 20b and SIGNAL 20b.

$\langle \text{Unify whitespace and comments 48a} \rangle \equiv$

This code is used in chunk 21.

$$\langle \textit{Tokenize numbers 48b} \rangle \equiv$$

Uses SIGNAL 20b.

This code is used in chunk 25b.

6.4 Strings and characters

49a *⟨Check and mask the strings 49a⟩*≡

$$\begin{aligned} &0 \neq \#lin \leftarrow \underline{1} \triangleright \phi \text{ msk} \leftarrow \# \backslash \text{''''} = IN \circ I \text{'' pos} : \{ \\ &EM \leftarrow \text{'SYNTAX ERROR: UNBALANCED STRING', ('S' / \text{''} 2 \leq \#lin), CR} \\ &EM, \leftarrow \text{quotelines } \epsilon(\text{msk} / \text{'' pos})[lin] \\ &EM \sqsubseteq \text{SIGNAL } 2 \\ &\} \emptyset \end{aligned}$$

This code is used in chunk 21.
 Uses `quotelines` 20b and `SIGNAL` 20b.

49b *⟨Tokenize strings 49b⟩*≡

$$\begin{aligned} &\text{end} \leftarrow 1 + \text{pos} \diamond t[\underline{i} \leftarrow \underline{1} 2 < / 0; \text{msk}] \leftarrow C \diamond \text{end}[i] \leftarrow \text{end}[\underline{1} 2 > / \text{msk}; 0] \\ &t \text{ pos } \text{end} / \text{''} \leftarrow (t \neq 0) \vee \sim \text{msk} \end{aligned}$$

This code is used in chunk 21.

6.5 Variables

49c *⟨Tokenize variables 49c⟩*≡

$$\begin{aligned} &t[\underline{i} \leftarrow \underline{1} 2 < / 0; \text{vm} \leftarrow (\sim \text{dm}) \wedge x \in \text{alp, num}] \leftarrow V \diamond \text{end}[i] \leftarrow \text{end} / \text{''} 2 > / \text{vm}; 0 \\ \\ &\text{A Tokenize } \alpha, \omega \text{ formals} \\ &\text{fm} \leftarrow \{ \text{mm} \leftarrow \phi \triangleright (> \circ \triangleright, \vdash) / \phi \text{ m} \leftarrow \alpha = ' ', \omega \diamond 1 \downarrow \text{''} (\text{mm} \wedge \sim \text{m} 1) (\text{mm} \wedge \text{m} 1 \leftarrow 1 \phi \text{ m}) \} \\ &\text{am } \text{aam} \leftarrow \alpha' \text{ fm } x \diamond \text{wm } \text{wvm} \leftarrow \omega' \text{ fm } x \\ &((\text{am} \vee \text{wm}) / t) \leftarrow A \diamond ((\text{aam} \vee \text{wvm}) / t) \leftarrow P \diamond ((\text{aam} \vee \text{wvm}) / \text{end}) \leftarrow \text{end} / \text{''} 1 \phi \text{ aam} \vee \text{wvm} \end{aligned}$$

This code is used in chunk 21.

49d *⟨Check for out of context dfns formals 49d⟩*≡

$$\vee / (d = 0) \wedge (t = P) \wedge IN[\text{pos}] \epsilon \alpha \omega : \text{'DFN FORMAL REFERENCED OUTSIDE DFNS'} \sqsubseteq \text{SIGNAL } 2$$

 This code is used in chunk 21.
 Uses `SIGNAL` 20b.

49e *⟨Convert α and ω to V nodes 49e⟩*≡

$$t \leftarrow V @ (i \leftarrow \underline{1} (t = A) \wedge n \epsilon \text{''} \alpha \omega') \vdash t \diamond \text{vb}[i] \leftarrow i$$

 This code is used in chunk 22.

49f *⟨Convert $\alpha\alpha$ and $\omega\omega$ to P2 nodes 49f⟩*≡

$$k[\underline{1} (t = P) \wedge n \epsilon \alpha \alpha \text{' ' } \omega \omega'] \leftarrow 2$$

 This code is used in chunk 22.

49g *⟨Node \leftrightarrow Generator mapping 24c⟩*+≡

$$\begin{aligned} &\text{gck}, \leftarrow (V \ 0) (V \ 1) (V \ 2) (V \ 3) (V \ 4) \\ &\text{gcv}, \leftarrow \text{'Va' 'Va' 'Vf' 'Vo' 'Vo'} \end{aligned}$$

 This code is used in chunk 25b.

50a $\langle \text{Node-specific code generators 24e} \rangle + \equiv$

$$\begin{aligned} & \text{Va} \leftarrow \{ \text{id} \leftarrow (|4 \rightarrow \alpha) \rightarrow ' ' 'r' 'l' 'aa' 'ww', 5 \downarrow \text{sym} \\ & \quad \text{z} \leftarrow \text{'*stkhdt++ = retain_cell(' , id, ')}; ' \\ & \quad \text{z} \} \end{aligned}$$

This code is used in chunk 25b.
 Uses `retain_cell` 40a.

6.6 Arrays

50b $\langle \text{Mark atoms, characters, and numbers as kind 1 50b} \rangle \equiv$

$$\text{k}[\underline{\text{t}} \in \text{A C N}] \leftarrow 1$$

This code is used in chunk 22.

50c $\langle \text{Strand arrays into atoms 50c} \rangle \equiv$

$$\begin{aligned} & \text{i} \leftarrow | \text{i} \rightarrow \text{km} \rightarrow 0 < \text{i} \rightarrow \text{i}[\text{A} | (\text{i}, \ddot{\text{z}} \leftarrow \text{up}[\text{i}]), \text{p}[\text{i} \rightarrow \underline{\text{t}}[\text{p}] \in \text{B Z}]] \\ & \text{msk} \leftarrow (\text{t}[\text{i}] \in \text{C N}) \vee \text{msk} \wedge \rightarrow 1 \quad \neg 1 \vee . \phi \leftarrow \text{msk} \leftarrow \text{km} \wedge (\text{t}[\text{i}] \in \text{A C N V Z}) \wedge \text{k}[\text{i}] = 1 \\ & \text{np} \leftarrow (\neq \text{p}) + \text{i} \neq \text{ai} \rightarrow \text{i} \neq \text{am} \rightarrow 2 \neq \text{msk}; 0 \quad \diamond \text{p} \leftarrow (\text{np} @ \text{ai} \neq \text{p})[\text{p}] \quad \diamond \text{p}, \leftarrow \text{ai} \quad \diamond \text{km} \rightarrow 2 < \neq 0; \text{msk} \\ & \text{t k n pos end}(\neg, \text{I}) \leftarrow \text{c ai} \quad \diamond \text{k}[\text{ai}] \leftarrow 1 \quad 6[\vee \neq \text{msk} \leq \text{t}[\text{i}] \neq \text{N}] \\ & \text{t n pos}(\neg @ \text{ai} \ddot{\text{z}}) \leftarrow \text{A}(\text{c} ' ') (\text{pos}[\text{km} \neq \text{i}]) \quad \diamond \text{p}[\text{msk} \neq \text{i}] \leftarrow \text{ai}[(\text{msk} \leftarrow \text{msk} \wedge \sim \text{am}) \neq 1 + \neq \text{km}] \\ & \text{i} \rightarrow \underline{\text{t}}(\text{t}[\text{p}] = \text{A}) \wedge (\text{k}[\text{p}] = 6) \wedge \text{t} = \text{N} \\ & \text{p}, \leftarrow \text{i} \quad \diamond \text{t k n pos end}(\neg, \text{I}) \leftarrow \text{c i} \quad \diamond \text{t k n}(\neg @ \text{i} \ddot{\text{z}}) \leftarrow \text{A 1}(\text{c} ' ') \end{aligned}$$

This code is used in chunk 22.

50d $\langle \text{Count strand and indexing children 50d} \rangle \equiv$

$$\text{n}[\underline{\text{t}}(\text{t} \in \text{A E}) \wedge \text{k} = 6] \leftarrow 0 \quad \diamond \text{n}[\text{p} \neq \ddot{\text{z}}(\text{t}[\text{p}] \in \text{A E}) \wedge \text{k}[\text{p}] = 6] \leftarrow 1$$

This code is used in chunk 23.

50e $\langle \text{Node} \leftrightarrow \text{Generator mapping 24c} \rangle + \equiv$

$$\begin{aligned} & \text{gck}, \leftarrow (\text{A 1})(\text{A 6}) \\ & \text{gcv}, \leftarrow \text{'Aa' 'As'} \end{aligned}$$

This code is used in chunk 25b.

50f $\langle \text{Declare top-level array structures 50f} \rangle \equiv$

$$\begin{aligned} & \text{k}[\omega] = 1 : \{ \\ & \quad \text{z} \leftarrow \text{'struct array *', n, '}; ' \\ & \quad \text{z} \} \omega \end{aligned}$$

This code is used in chunk 25a.

6.7 Primitives

50g $\langle \text{Node} \leftrightarrow \text{Generator mapping 24c} \rangle + \equiv$

$$\begin{aligned} & \text{gck}, \leftarrow (\text{P 0})(\text{P 1})(\text{P 2})(\text{P 3})(\text{P 4}) \\ & \text{gcv}, \leftarrow \text{'Pv' 'Pv' 'Pf' 'Po' 'Po'} \end{aligned}$$

This code is used in chunk 25b.

51a *⟨Node-specific code generators 24e⟩*≡
`Pf←{id←(syms⊔sym[|4>α])>nams
 z←c'*stkhd++ = retain_cell(' ,id, ');'
 z}`

This code is used in chunk 25b.
 Uses `retain_cell` 40a.

6.7.1 APL Primitives

51b *⟨Tokenize primitives and atoms 51b⟩*≡
`t[⊂(∼dm)∧x∈prms]←P ⋄ t[⊂x∈syna]←A`

This code is used in chunk 21.

51c *⟨Mark APL primitives with appropriate kinds 51c⟩*≡
`k[⊂ne, "prmfs"]←2 ⋄ k[⊂ne, "prmmo"]←3 ⋄ k[⊂ne, "prmdo"]←4
 k[⊂ne, "prmfo"]←5
 k[i←⊂msk←(n∈c, 'o')∧1ϕn∈c, '.']←3 ⋄ end[i]←end[i+1] ⋄ n[i]←c, 'o.'
 t k n pos end/⊂←msk←∼1ϕmsk ⋄ p←(⊂∼msk)(⊂-1+⊂)msk/⊂`

This code is used in chunk 22.

6.7.2 System Functions and Variables

51d *⟨Tokenize system variables 51d⟩*≡
`si←⊂('□'=IN[pos])∧1ϕt=V
 t[si]←S ⋄ end[si]←end[si+1] ⋄ t[si+1]←0`

This code is used in chunk 21.

52a *⟨Verify that system variables are defined 52a⟩≡*

```

SYSV←,,"Á" 'A' 'AI' 'AN' 'AV' 'AVU' 'BASE' 'CT' 'D' 'DCT' 'DIV' 'DM'
SYSV,←,,"DMX" 'EXCEPTION' 'FAVAIL' 'FNAMES' 'FNUMS' 'FR' 'IO' 'LC' 'LX'
SYSV,←,,"ML" 'NNAMES' 'NNUMS' 'NSI' 'NULL' 'PATH' 'PP' 'PW' 'RL' 'RSI'
SYSV,←,,"RTL" 'SD' 'SE' 'SI' 'SM' 'STACK' 'TC' 'THIS' 'TID' 'TNAME' 'TNUMS'
SYSV,←,,"TPOOL" 'TRACE' 'TRAP' 'TS' 'USING' 'WA' 'WSID' 'WX' 'XSI'
SYSF,←,,"ARBIN" 'ARBOU' 'AT' 'C' 'CLASS' 'CLEAR' 'CMD' 'CONV' 'CR' 'CS' 'CSV'
SYSF,←,,"CY" 'DF' 'DL' 'DQ' 'DR' 'DT' 'ED' 'EM' 'EN' 'EX' 'EXPORT'
SYSF,←,,"FAPPEND" 'FCHK' 'FCOPY' 'FCREATE' 'FDROP' 'FERASE' 'FFT' 'IFFT'
SYSF,←,,"FHIST" 'FHOLD' 'FIX' 'FLIB' 'FMT' 'FPROPS' 'FRDAC' 'FRDCI' 'FREAD'
SYSF,←,,"FRENAME' 'FREPLACE' 'FRESIZE' 'FSIZE' 'FSTAC' 'FSTIE' 'FTIE'
SYSF,←,,"FUNTIE' 'FX' 'INSTANCES' 'JSON' 'KL' 'LOAD' 'LOCK' 'MAP' 'MKDIR'
SYSF,←,,"MONITOR' 'NA' 'NAPPEND' 'NC' 'NCOPY' 'NCREATE' 'NDELETE' 'NERASE'
SYSF,←,,"NEW' 'NEXISTS' 'NGET' 'NINFO' 'NL' 'NLOCK' 'NMOVE' 'NPARTS'
SYSF,←,,"NPUT' 'NQ' 'NR' 'NREAD' 'NRENAME' 'NREPLACE' 'NRESIZE' 'NS'
SYSF,←,,"NSIZE' 'NTIE' 'NUNTIE' 'NXLATE' 'OFF' 'OR' 'PFKEY' 'PROFILE'
SYSF,←,,"REFS' 'SAVE' 'SH' 'SHADOW' 'SIGNAL' 'SIZE' 'SR' 'SRC' 'STATE'
SYSF,←,,"STOP' 'SVC' 'SVO' 'SVQ' 'SVR' 'SVS' 'TCNUMS' 'TGET' 'TKILL' 'TPUT'
SYSF,←,,"TREQ' 'TSYNC' 'UCS' 'VR' 'VFI' 'WC' 'WG' 'WN' 'WS' 'XML' 'XT'
SYSD←,,"OPT' 'R' 'S'
v/mask←(t=S)∧n∈'□',,"SYSV,SYSF,SYSD:{
ERR←2'INVALID SYSTEM VARIABLE, FUNCTION, OR OPERATOR'
ERR SIGNAL←pos[ω]{α+ιω-α}"end[ω]
}lmsk

```

This code is used in chunk 22.
Uses SIGNAL 20b.

52b *⟨Mark system variables as P nodes with appropriate kinds 52b⟩≡*

```

k[l(t=S)∧n∈'□',,"SYSV]←1 ♦ k[l(t=S)∧n∈'□',,"SYSF]←2 ♦ k[l(t=S)∧n∈'□',,"SYSD]←4
t[l t=S]←P

```

This code is used in chunk 22.

6.8 Brackets

6.8.1 Indexing

52c *⟨Convert ; groups within brackets into Z nodes 52c⟩≡*

```

_←p[i]{k[z←;g←g←ωc-1φIN[pos[ω]]∈';']]+1 ♦ t[z]←Z P[1≠'g]}i←lt[p]=-1

```

This code is used in chunk 22.

52d *⟨Verify brackets have function/array target 52d⟩≡*

```

x←{ω≠~∧t[ω]=-1}Uφ"x
0v.≠"x:'BRACKET SYNTAX REQUIRES FUNCTION OR ARRAY TO ITS LEFT'□SIGNAL 2

```

This code is used in chunk 54b.
Uses SIGNAL 20b.

- 53a *⟨Enclose $V[X; \dots]$ for expression parsing 53a⟩*≡
 $i \leftarrow i[\Delta p[i \leftarrow \underline{1}(t[p] \in B \ Z) \wedge (k[p] = 1) \wedge p \neq i \neq p]] \diamond j \leftarrow i \neq j m \leftarrow t[i] = -1$
 $t[j] \leftarrow A \diamond k[j] \leftarrow -1 \diamond p[i \neq 1 \phi j m] \leftarrow j$
 This code is used in chunk 22.
- 53b *⟨Rationalize $V[X; \dots]$ 53b⟩*≡
 $i \leftarrow i[\Delta p[i \leftarrow \underline{1}(t[p] = A) \wedge k[p] = -1]] \diamond msk \leftarrow -2 \neq -1, ip \leftarrow p[i] \diamond ip \leftarrow uip \diamond nc \leftarrow 2 \times \neq ip$
 $t[ip] \leftarrow E \diamond k[ip] \leftarrow 2 \diamond n[ip] \leftarrow c \text{' ' } \diamond p[msk \neq i] \leftarrow msk \neq (\neq p) + 1 + 2 \times -1 + + \backslash \sim msk$
 $p, \neq 2 \neq ip \diamond t, \neq nc p \text{ P } E \diamond k, \neq nc p 2 \ 6 \diamond n, \neq nc p, \text{' ' ' ' ' '}$
 $pos, \neq 2 \neq pos[ip] \diamond end, \neq \epsilon(1 + pos[ip]), \neq end[ip] \diamond pos[ip] \leftarrow pos[i \neq \sim msk]$
 This code is used in chunk 22.
- 53c *⟨Symbol \leftrightarrow Name mapping 24b⟩*+≡
 $syms, \leftarrow c, \text{' ; ' } \diamond nams, \leftarrow c \text{' span '}$
 This code is used in chunk 25b.
- 53d *⟨Node \leftrightarrow Generator mapping 24c⟩*+≡
 $gck, \leftarrow c \text{ E } 6$
 $gcv, \leftarrow c \text{ ' E i '}$
 This code is used in chunk 25b.
- ## 6.8.2 Axis Operator
- 53e *⟨Rationalize $F[X]$ syntax 53e⟩*≡
 $_ \leftarrow p[i] \{$
 $\triangleright m \leftarrow t[\omega] = -1 : \text{'SYNTAX ERROR: NOTHING TO INDEX'} \square \text{SIGNAL } 2$
 $k[\omega \neq \sim m \wedge -1 \phi (k[\omega] \in 2 \ 3 \ 5) \vee -1 \phi k[\omega] = 4] \leftarrow 4$
 $0 \} \exists i \leftarrow \underline{1}(t[p] \in B \ Z) \wedge (p \neq i \neq p) \wedge k[p] \in 1 \ 2$
 $i \leftarrow \underline{1}(t = -1) \wedge k = 4 \diamond j \leftarrow \underline{1}(t[p] = -1) \wedge k[p] = 4$
 $(\neq i) \neq j : \{$
 $2 \text{'AXIS REQUIRES SINGLE AXIS EXPRESSION' SIGNAL } \epsilon pos[\omega] + i \text{' ' end}[\omega] - pos[\omega]$
 $\} \triangleright, \neq \{ c \alpha \neq \sim 1 < \neq \omega \} \exists p[j]$
 $\vee \neq msk \leftarrow t[j] \neq Z : \{$
 $2 \text{'AXIS REQUIRES NON-EMPTY AXIS EXPRESSION' SIGNAL } \epsilon pos[\omega] + i \text{' ' end}[\omega] - pos[\omega]$
 $\} msk \neq p[j]$
 $p[j] \leftarrow p[i] \diamond t[i] \leftarrow P \diamond end[i] \leftarrow 1 + pos[i]$
 This code is used in chunk 22.
 Uses SIGNAL 20b.

6.9 Bindings and Types

54a *Parse Binding nodes 54a* \equiv
 A Mark bindable nodes
 $bm \leftarrow (t=V) \vee (t=A) \wedge n \in \text{' '}$
 $bm \leftarrow \{bm \rightarrow p[i] \mid \{bm[\alpha] \leftarrow (V^{-1} \equiv t[\omega]) \vee \wedge \neg bm[\omega]\} \exists i \leftarrow \underline{1} (\sim bm[p]) \wedge t[p]=Z\} * \equiv bm$

A Binding nodes
 $\underline{p}[i] \{$
 $t[\omega] \leftarrow (n[\omega] \in c, ' \leftarrow ') \wedge 0, \neg 1 \downarrow bm[\omega] \} \leftarrow B$
 $b \vee \leftarrow \{(\supset x)(1 \downarrow x \leftarrow \omega \neg \{t[\supset \omega]=B\} \neg \omega)\} \neg 1 \phi \neg \omega \neg 1, \neg 1 \downarrow t[\omega] \in P \ B$
 $\vee \neg \sim bm[\epsilon v]: \text{'CANNOT BIND ASSIGNMENT VALUE'}$ SIGNAL 2
 $p[\omega] \leftarrow (\alpha, b)[0, \neg 1 \downarrow \neg t[\omega]=B]$
 $n[b] \leftarrow n[\epsilon v] \diamond t[\epsilon v] \leftarrow 7 \diamond pos[b] \leftarrow pos[\epsilon v] \diamond end[b] \leftarrow end[\supset \phi \omega]$
 $0\} \exists i \leftarrow \underline{1} (t[p]=Z) \wedge p \neq i \neq p$
 $t \ k \ n \ pos \ end \neg \neg \neg msk \leftarrow t \neq 7 \diamond p \leftarrow (\underline{1} \sim msk) (\neg 1 + \underline{1}) msk \neg p$

This code is used in chunk 22.
 Uses SIGNAL 20b.

54b *Infer the type of bindings, groups, and variables 54b* \equiv
 $z \ x \neg \downarrow \phi p[i] \{ \alpha \omega \} \exists i \leftarrow \underline{1} (t[p] \in B \ Z) \wedge p \neq i \neq p$
Verify brackets have function/array target 52d
 $\underline{p} \leftarrow \{$
 $k[msk \neg z] \leftarrow k[x \neg \neg msk \leftarrow (k[\supset x] \neq 0) \wedge 1 = \neg x]$
 $z \ x \neg \neg \neg \neg msk$

 $k[z \neg \neg msk \leftarrow k[\supset x]=4] \leftarrow 3$
 $z \ x \neg \neg \neg \neg msk$

 $k[z \neg \neg msk \leftarrow \{(2 \ 3 \ 5 \neg k[\supset \omega]) \vee 4 = (\omega, \neq k)[0 \neg \neg \neg k[\omega]=1] \neg k, 0\} \circ \phi x] \leftarrow 2$
 $z \ x \neg \neg \neg \neg msk$

 $k[z \neg \neg msk \leftarrow k[\supset \phi x]=1] \leftarrow 1$
 $z \ x \neg \neg \neg \neg msk$

 $k[i] \leftarrow k[vb[i \leftarrow \underline{1} t=V]]$
 $\neq z \} * (=v0=\neg) \neq z$
 'FAILED TO INFER ALL BINDING TYPES' assert 0 $\neq z$:

This code is used in chunk 22.

54c *Parse dyadic operator bindings 54c* \equiv
 A PARSE B $\leftarrow D \dots$
 A PARSE B $\leftarrow \dots D$

This code is used in chunk 22.

This code is used in chunk 25b.

This code is used in chunk 25b.
Uses `retain_cell` 40a.

6.10 Assignments

This code is used in chunk 22.

- 56a $\langle \textit{Symbol} \leftrightarrow \textit{Name mapping} \text{ 24b} \rangle + \equiv$
`syms, ← c, '←' ◇ nams, ← c 'get'`
 This code is used in chunk 25b.
- 56b $\langle \textit{Node} \leftrightarrow \textit{Generator mapping} \text{ 24c} \rangle + \equiv$
`gck, ← c E 4`
`gcv, ← c 'Eb'`
 This code is used in chunk 25b.
- 56c $\langle \textit{Cell type names} \text{ 36c} \rangle + \equiv$
`, CELL_BOX`
 This code is used in chunk 36b.
 Defines:
`CELL_BOX`, used in chunk 57.
- 56d $\langle \textit{C runtime structures} \text{ 36a} \rangle + \equiv$
`struct cell_box {`
 $\langle \textit{Common cell fields} \text{ 35} \rangle$
`void *value;`
`};`
 This code is used in chunk 33.
 Defines:
`box`, used in chunks 57a, 58b, and 80.

57a *⟨Box definitions 57a⟩*≡

```

DECLSPEC int
mk_box(struct cell_box **box, void *value)
{
    *box = malloc(sizeof(struct cell_box));

    if (*box == NULL)
        return 1;

    (*box)->ctyp    = CELL_BOX;
    (*box)->refc    = 1;
    (*box)->value   = value;

    return 0;
}

DECLSPEC void
release_box(struct cell_box *box)
{
    if (box == NULL)
        return;

    box->refc--;

    if (box->refc)
        return;

    release_cell(box->value);
    free(box);
}

```

This code is used in chunk 58a.

Defines:

mk_box, used in chunk 57b.

release_box, used in chunk 57.

Uses box 56d, CELL_BOX 56c, ctyp 35, DECLSPEC 34b, refc 35, and release_cell 39a.

57b *⟨C runtime declarations 38a⟩*+≡

```

DECLSPEC int mk_box(struct cell_box **, void *);
DECLSPEC void release_box(struct cell_box *);

```

This code is used in chunk 33.

Uses DECLSPEC 34b, mk_box 57a, and release_box 57a.

57c *⟨Cell release cases 39c⟩*+≡

```

case CELL_BOX:
    release_box(cell);
    break;

```

This code is used in chunk 39a.

Uses CELL_BOX 56c and release_box 57a.

58a $\langle box.c\ 58a \rangle \equiv$
`#include <stdlib.h>`

`#include "codfns.h"`

(Box definitions 57a)

Root chunk (not used in this document).

Defines:

`box.c`, used in chunk 58b.

Uses `codfns 7` and `codfns.h 33`.

58b $\langle Tangle\ Commands\ 8 \rangle_+ \equiv$
`echo "Tangling rtm/box.c..."`
`notangle -R'box.c' codfns.nw > rtm/box.c`

This code is used in chunk 83.

Uses `box 56d`, `box.c 58a`, and `codfns 7`.

6.11 Expressions

58c $\langle Parse\ brackets\ and\ parentheses\ into\ ^{-1}\ and\ \mathbb{Z}\ nodes\ 58c \rangle \equiv$
`_←p[i]{`
`x←IN[pos[ω]]`
`bd←+λbm←(bo←'('=x)+-bc←')'=x`
`pd←+λpm←(po←'('=x)+-pc←')'=x`
`0≠φbd:{`
`ix←pos[ω]{x+ι([≠ω)-x←[≠α]}ö{ω≠0≠bd}end[ω]`
`2'UNBALANCED BRACKETS'SIGNAL ix`
`}ω`
`0≠φpd:{`
`ix←pos[ω]{x+ι([≠ω)-x←[≠α]}ö{ω≠0≠pd}end[ω]`
`2'UNBALANCED PARENTHESES'SIGNAL ix`
`}ω`
`(po≠bd)∨.≠φpc≠bd:{`
`'OVERLAPPING BRACKETS AND PARENTHESES'□SIGNAL 2`
`}ω`
`p[ω]←(α,ω)[1+^{-1}@{ω=ι≠ω}D2P +λ^{-1}φbm+pm]`
`t[bo≠ω]←^{-1} ♦ t[po≠ω]←Z`
`end[po≠ω]←end[φpc≠ω] ♦ end[bo≠ω]←end[φbc≠ω]`
`0}∃i←_1(t[p]=Z)∧p≠ι≠p`
`t k n pos end≠↖←msk←IN[pos]ε')' ♦ p←(_1~msk)(ι-1+_1)msk≠p`

This code is used in chunk 22.

Uses `SIGNAL 20b`.

58d $\langle Group\ function\ and\ value\ expressions\ 58d \rangle \equiv$
`i km←≠≠p[i]{(α;ω)(0,1∨ω)}∃i←_1(t[p]∈B Z)∧(p≠ι≠p)∧k[p]∈1 2`

This code is used in chunk 22.

59a $\langle \text{Lift and flatten expressions } 59a \rangle \equiv$
 $p[i] \leftarrow p[x \leftarrow p \text{ I}@\{\sim t[p[\omega]] \in F \ G\} \ddot{x} \equiv i \leftarrow \underline{1} t \in G \ A \ B \ C \ E \ O \ P \ V] \diamond j \leftarrow (\phi i)[\Delta \phi x]$
 $p \ t \ k \ n \ r\{\alpha[\omega]@i \vdash \alpha\} \leftarrow c \ j \diamond p \leftarrow (i@j \vdash \neg p)[p]$
 This code is used in chunk 23.

6.11.1 Value Expressions

59b $\langle \text{Parse value expressions } 59b \rangle \equiv$
 $i \leftarrow km \leftarrow \neg p[i]\{(\alpha; \omega)(0, (2 \leq \omega) \wedge 1 \vee \omega)\} \boxplus i \leftarrow \underline{1} (t[p] \in B \ Z) \wedge (k[p]=1) \wedge p \neq \neg p$
 $msk \leftarrow m2 \vee fm \wedge \sim 1 \phi m2 \leftarrow km \wedge (1 \phi km) \wedge \sim fm \leftarrow (t[i]=0) \vee (t[i] \neq A) \wedge k[i]=2$
 $t, \leftarrow E p \ddot{\sim} xc \leftarrow \neg msk \diamond k, \leftarrow msk \neg msk + m2 \diamond n, \leftarrow xc p c '$
 $pos, \leftarrow pos[msk \neg i] \diamond end, \leftarrow end[p[msk \neg i]]$
 $p, \leftarrow msk \neg 1 \phi (i \times \sim km) + km \times x \leftarrow \neg 1 + (\neg p) ++ \neg msk \diamond p[km \neg i] \leftarrow km \neg x$
 This code is used in chunk 22.

59c $\langle \text{Node} \leftrightarrow \text{Generator mapping } 24c \rangle + \equiv$
 $gck, \leftarrow (E \ 1)(E \ 2)$
 $gcv, \leftarrow 'Em' \ 'Ed'$
 This code is used in chunk 25b.

59d $\langle \text{Node-specific code generators } 24e \rangle + \equiv$
 $Em \leftarrow \{$
 $z \leftarrow c 'c = *--stkhd; '$
 $z, \leftarrow c 'w = *--stkhd; '$
 $z, \leftarrow c '(c \rightarrow fn)((struct \ array \ **)stkhd++, \ NULL, \ w, \ c \rightarrow fv); '$
 $z, \leftarrow c 'release_cell(c); '$
 $z, \leftarrow c 'release_cell(w); '$
 $z \}$
 This code is used in chunk 25b.

6.11.2 Function Expressions

60a *⟨Parse function expressions 60a⟩*≡

```

A Mask and verify dyadic operator right operands
(dm←¬1φ(k[i]=4)∧t[i]∈F P V Z)∨.∧(¬km)∨k[i]∈0 3 4:{
'MISSING RIGHT OPERAND'␣SIGNAL 2
}θ

A Refine schizophrenic types
k[i]≠(k[i]=5)∧dm∨¬1φ(¬km)∨(¬dm)∧k[i]∈1 6]←2 ⋄ k[i]≠k[i]=5]←3

A Rationalize °.
jm←(t[i]=P)∧n[i]∈c, '°.'
jmv.∧1φ(¬km)∨k[i]∈3 4:'MISSING OPERAND TO °.'␣SIGNAL 2
p←((ji+jm)÷i)@((jj+i)÷¬1φjm)⌊p[p] ⋄ t[ji,jj]←t[jj,ji] ⋄ k[ji,jj]←k[jj,ji]
n[ji,jj]←n[jj,ji] ⋄ pos[ji,jj]←pos[ji,ji] ⋄ end[ji,jj]←end[jj,jj]

A Mask and verify monadic and dyadic operator left operands
∨≠msk←(dm¬2φ¬km)∨(¬1φ¬km)∧mm←(k[i]=3)∧t[i]∈F P V Z:{
2'MISSING LEFT OPERAND'SIGNAL εpos[ω]+⌊end[ω]-pos[ω]
}i≠msk
msk←dm∨mm

A Parse function expressions
np←(≠p)+⌊xc≠oi←msk÷i ⋄ p←(np@oi⌊p)[p] ⋄ p,←oi ⋄ t k n pos end(¬,I)←coi
p[g÷i]←oi[(g←(¬msk)∧(1φmsk)∨2φdm)÷xc-φ+∧φmsk]
p[g÷oi]←(g←msk÷(1φmm)∨2φdm)÷1φoi ⋄ t[oi]←0 ⋄ n[oi]←c'
pos[oi]←pos[g÷i][msk÷¬1+∧g←(¬msk)∧(1φmm)∨2φdm]
ol←1+(k[i]≠(2φmm)∨3φdm)=4)∨k[i]≠(1φmm)∨2φdm]∈2 3
or←(msk÷dm)∧1+k[dm÷i]=2
k[oi]←3 3⌊tor ol

This code is used in chunk 22.
Uses SIGNAL 20b.

```

60b *⟨Node ↔ Generator mapping 24c⟩*+≡

```

gck,←(0 1)(0 2)(0 4) (0 5) (0 7) (0 8)
gcv,←'Ov' 'Of' 'Ovv' 'Ofv' 'Ovf' 'Off'

This code is used in chunk 25b.

```

6.12 Trains

60c *⟨Parse trains 60c⟩*≡

```

A TRAINS

This code is used in chunk 22.

```

6.13 Functions

61a *⟨Declare top-level function bindings 61a⟩*≡
`k[ω] ← 0 2 : {
 z ← c 'int'
 z, ← c n, '(struct array **z, struct array *l, struct array *r, void *fv[]);'
 z, ← c '
 z } ω`

This code is used in chunk 24e.

61b *⟨Declare top-level closures 61b⟩*≡
`k[ω] ← 2 : {
 z ← c 'struct closure *', n, ';' ;
 z, ← c '
 (DWA Function Export 29)
 z } ω`

This code is used in chunk 25a.

6.13.1 D-fns

61c *⟨Compute dfns regions and type, with } as a child 61c⟩*≡
`t[⊥] { ' = x } ← F ◊ 0 ≠ d ← 1 φ + λ 1 1 0 [' { } ' l x] : ' UNBALANCED DFNS ' ⊔ SIGNAL 2`
 This code is used in chunk 21.
 Uses SIGNAL 20b.

61d *⟨Compute the nameclass of dfns 61d⟩*≡
`k ← 2 × t ∈ F ◊ k[upf ⌘ (t = P) ∧ n ∈ c 'αα'] ← 3 ◊ k[upf ⌘ (t = P) ∧ n ∈ c 'ωω'] ← 4`
 This code is used in chunk 22.

61e *⟨Wrap all dfns expression bodies as Z nodes 61e⟩*≡
`⊥ ← p[i] { end[α] ← end[φω] ◊ gz ⌘ ω c ⌘ 1, 1 ⊥ t[ω] = Z } ⊔ i ← ⊥ t[p] = F
 'Non-Z dfns body node' assert t[⊥ t[p] = F] = Z :`
 This code is used in chunk 22.

61f *⟨Anchor variables to earliest binding in the matching frame 61f⟩*≡
`rf ← 1 ⊔ { ~ t[ω] ∈ F G M } p[rz ← I ⊔ { ~ (t[ω] = Z) ∧ (t[p[ω]] ∈ F G M) ∨ p[ω] = ω } * ⌘ p]
 rf[i] ← p[i ← ⊥ t = G] ◊ rz[i] ← i ◊ rf ← rf I ⊔ { rz ∈ p[i] ⊢ ◊ ⊔ i ← ⊥ t[p] = G } rf
 mk ← { α[ω], ; n[ω] }
 fr ← rf mk ⊢ fb ← fb[i ⌘ rf mk ⊢ fb ← fb I ◊ (i ⌘) U ⊔ rz mk ⊢ fb ← ⊥ t = B] ◊ fb, ← 1
 vb ← fb[fr ⊢ rf mk i] ⊔ (i ← ⊥ t = V) ⊢ 1 p ⌘ p
 vb[i ⌘ (rz[i] < rz[b]) ∨ (rz[i] = rz[b]) ∧ i ≥ b + vb[i + i ⌘ vb[i] ≠ 1]] ← 1
 ⊥ ← { z / ⌘ 1 = vb[1] z] + fb[fr ⊢ n I ⊔ 1 ⊢ z + rf I ⊔ 0 ⊢ ω] } * ⌘ { rf[ω], ; ω } ⊥ (t = V) ∧ vb = 1
 √ msk ← (t = V) ∧ vb = 1 : {
 6 'ALL VARIABLES MUST REFERENCE A BINDING' SIGNAL ⊢ pos[ω] { α + i ω - α } ⌘ end[ω]
 } ⊥ msk`

This code is used in chunk 22.

62a *⟨Lift dfns to the top-level 62a⟩*≡
 $p, \leftarrow n[i] \leftarrow (\neq p) + i \neq i \leftarrow \underline{1} (t = F) \wedge p \neq i \neq p \diamond t \text{ k n r}(\neg, I) \leftarrow c i \diamond p \text{ r } I \leftarrow c n[i] @ i \neg i \neq p$
 $t[i] \leftarrow C$

This code is used in chunk 23.

62b *⟨Wrap expressions as binding or return statements 62b⟩*≡
 $i \leftarrow (\underline{1} (\neg t \in F \text{ G}) \wedge t[p] = F), \{\omega \neq \omega \mid i \neq \omega\} \underline{1} t[p] = G \diamond p \text{ t k n r} \neq c m \leftarrow 2 @ i \neg 1 p \neq p$
 $p \text{ r } i \text{ I} \neq c j \leftarrow (+ \backslash m) - 1 \diamond n \leftarrow j \text{ I} @ (0 \leq \neg) n \diamond p[i] \leftarrow j \neg i - 1$
 $k[j] \leftarrow (k[r[j]] = 0) \vee 0 @ (\{ \triangleright \phi \omega \} \exists p[j]) \neg (t[j] = B) \vee (t[j] = E) \wedge k[j] = 4 \diamond t[j] \leftarrow E$

This code is used in chunk 23.

62c *⟨Node ↔ Generator mapping 24c⟩*+≡
 $gck, \leftarrow (E \neg 1)(E \text{ 0})$
 $gcv, \leftarrow 'Ek' \neg 'Er'$

This code is used in chunk 25b.

62d *⟨Compute slots and frames 62d⟩*≡
 $\text{A Compute slots for each frame}$
 $s \leftarrow \neg 1, \neg \in i \neg n[ux] \leftarrow \neg \circ \neq \exists x \leftarrow 0 \exists qe \leftarrow u \text{ I} \circ \neq \neg r n \leftarrow r[b], \neg n[b \leftarrow \underline{1} t = B]$
 $\text{A Compute frame depths}$
 $d \leftarrow (\neq p) \uparrow d \diamond d[i \leftarrow \underline{1} t = F] \leftarrow 0 \diamond _ \leftarrow \{z \neg d[i] \leftarrow \omega \neq z \neg r[\omega]\} \neq i \diamond f \leftarrow d[0 \exists qe], \neg 1$

This code is used in chunk 23.

62e *⟨Symbol ↔ Name mapping 24b⟩*+≡
 $syms, \leftarrow c, ' \nabla ' \diamond nams, \leftarrow c 'this'$

This code is used in chunk 25b.

62f *⟨Node ↔ Generator mapping 24c⟩*+≡
 $gck, \leftarrow (C \text{ 1})(C \text{ 2})(F \text{ 2})(F \text{ 3})(F \text{ 4})$
 $gcv, \leftarrow 'Ca' \neg 'Cf' \neg 'Fn' \neg 'Fm' \neg 'Fd'$

This code is used in chunk 25b.

62g *⟨Node-specific code generators 24e⟩*+≡
 $Cf \leftarrow \{id \leftarrow \neq 4 \triangleright \alpha$
 $z \leftarrow c 'mk_closure((struct \text{ closure } **)stkhd++, fn', id, ', 0);'$
 $z\}$

This code is used in chunk 25b.

62h *⟨Node-specific code generators 24e⟩*+≡
 $E k \leftarrow \{$
 $z \leftarrow c 'release_cell(*--stkhd);'$
 $z, \leftarrow c ' '$
 $z\}$

This code is used in chunk 25b.

63a $\langle \text{Node-specific code generators 24e} \rangle + \equiv$

```

Er←{
  z ←c '*z = *--stkhd;'
  z,←c 'goto cleanup;'
  z,←c '
  z}

```

This code is used in chunk 25b.

63b $\langle \text{Node-specific code generators 24e} \rangle + \equiv$

```

Fn←{id←5▷α ◊ x←Q▷;fω ◊ t←2[]x ◊ k←3[]x
hsw←(t=0)∨(t=E)∧k∈1 2 ◊ hsa←((t=E)∧k=2)∨(t=0)∧k∈4 5 7 8
z ←c 'int'
z,←c 'fn',id,'(struct array **z, '
z,←c '      struct array *l, struct array *r, void *fv[])'
z,←c '{'
z,←c '          void      *stk[128];'
z,←c '          void      **stkhd;'
z,←hswf c ' void      *w;'
z,←hsaf c ' void      *a;'
z,←hswf c ' struct  closure *c;'
z,←c '
z,←c '          stkhd = &stk[0];'
z,←c '
z,← ' ',",,fdis"ω
z,←c '          *z = NULL;'
z,←c '
z,←c 'cleanup:'
z,←c '          return 0;'
z,←c '}'
z,←c '
z}

```

This code is used in chunk 25b.

6.13.2 Trad-fns

63c $\langle \text{Compute trad-fns regions 63c} \rangle \equiv$

```

v fZ≠t f 1φmsk←(d=0)∧'∇'=x: 'TRAD-FNS START/END LINES MUST BEGIN WITH ∇' □SIGNAL 2
0≠>tm←1φ≠λ(d=0)∧'∇'=x: 'UNBALANCED TRAD-FNS' □SIGNAL 2
v fZ≠t f 1 1∨.φ<(2>f tm);0: 'TRAD-FNS END LINE MUST CONTAIN ∇ ALONE' □SIGNAL 2

```

This code is used in chunk 21.

Uses SIGNAL 20b.

6.14 Guards

64a $\langle \text{Parse guards to } (G \ (Z \ \dots) \ (Z \ \dots)) \ 64a \rangle \equiv$

```

  _←p[i]{
    0=+7m←': '=IN[pos[ω]]:θ
    >m: 'EMPTY GUARD TEST EXPRESSION'□SIGNAL 2
    1<+7m: 'TOO MANY GUARDS'□SIGNAL 2
    t[α]←G ◇ p[ti←gz>tx cq←2↑(cθ);ω←1, -1↓m]←α ◇ k[ti]←1
    ci←p ◇ p, ←α ◇ t k pos end;←0 ◇ n, ←' ' ◇ k[gz cq, ci]←1
    0)∃i←1t[p[p]]=F

```

This code is used in chunk 22.
Uses SIGNAL 20b and TEST 16a.

64b $\langle \text{Lift guard tests } 64b \rangle \equiv$

```

  p[i]←p[x←1+i←{ω7~2|ι≠ω}]t[p]=G] ◇ t[i, x]←t[x, i] ◇ k[i, x]←k[x, i]
  n[x]←n[i] ◇ p←((x, i)@(i, x)ι≠p)[p]

```

This code is used in chunk 23.

64c $\langle \text{Node} \leftrightarrow \text{Generator mapping } 24c \rangle + \equiv$

```

  gck, ←cG 0
  gcv, ←c'Gd'

```

This code is used in chunk 25b.

6.14.1 Error Guards

6.15 Labels

64d $\langle \text{Identify label colons vs. others } 64d \rangle \equiv$

```

  t[1tm^(d=0)∧∈((~>)∧(<∧v∧))''' : '(t=Z)∈IN[pos]]←L

```

This code is used in chunk 21.

64e $\langle \text{Tokenize labels } 64e \rangle \equiv$

```

  ERR←'LABEL MUST CONSIST OF A SINGLE NAME'
  v7(Z≠t[li-1])∨(V≠t[li←1φmsk←t=L]):ERR □SIGNAL 2
  t[li]←L ◇ end[li]←end[li+1]
  d tm t pos end(7~)←c~msk

```

This code is used in chunk 21.
Uses SIGNAL 20b.

64f $\langle \text{Parse labels } 64f \rangle \equiv$

```

  A XXX: Parse labels

```

Root chunk (not used in this document).

6.16 Statements

6.16.1 What is a keyword?

65a *⟨Tokenize keywords 65a⟩*≡
 $ki \leftarrow \underline{1} (t=0) \wedge (d=0) \wedge (': '=IN[pos]) \wedge 1\phi t=V$
 $t[ki] \leftarrow K \diamond end[ki] \leftarrow end[ki+1] \diamond t[ki+1] \leftarrow 0$
 ERR←'EMPTY COLON IN NON-DFNS CONTEXT, EXPECTED LABEL OR KEYWORD'
 $\vee \neg (t=0) \wedge (d=0) \wedge ': '=IN[pos]: ERR \sqcap SIGNAL \ 2$

This code is used in chunk 21.

Uses SIGNAL 20b.

65b *⟨Check that all keywords are valid 65b⟩*≡
 $KW \leftarrow 'NAMESPACE' \ 'ENDNAMESPACE' \ 'END' \ 'IF' \ 'ELSEIF' \ 'ANDIF' \ 'ORIF' \ 'ENDIF'$
 $KW, \leftarrow 'WHILE' \ 'ENDWHILE' \ 'UNTIL' \ 'REPEAT' \ 'ENDREPEAT' \ 'LEAVE' \ 'FOR' \ 'ENDFOR'$
 $KW, \leftarrow 'IN' \ 'INEACH' \ 'SELECT' \ 'ENDSELECT' \ 'CASE' \ 'CASELIST' \ 'ELSE' \ 'WITH'$
 $KW, \leftarrow 'ENDWITH' \ 'HOLD' \ 'ENDHOLD' \ 'TRAP' \ 'ENDTRAP' \ 'GOTO' \ 'RETURN' \ 'CONTINUE'$
 $KW, \leftarrow 'SECTION' \ 'ENDSECTION' \ 'DISPOSABLE' \ 'ENDDISPOSABLE'$
 $KW, \leftarrow \ddot{\sim} \leftarrow ': '$
 $msk \leftarrow \sim KW \in \ddot{\sim} kws \leftarrow n \neg \ddot{\sim} km \leftarrow t = K$
 $\vee \neg msk: ('UNRECOGNIZED KEYWORD' \ , kws \supset \ddot{\sim} \supset \underline{1} msk) \sqcap SIGNAL \ 2$

This code is used in chunk 22.

Uses SIGNAL 20b.

6.16.2 Namespaces

65c *⟨Check that namespaces are at the top level 65c⟩*≡
 $msk \leftarrow kws \in ': NAMESPACE' \ '': ENDNAMESPACE'$
 $\vee \neg msk \wedge km \neg tm: 'NAMESPACE SCRIPTS MUST APPEAR AT THE TOP LEVEL' \sqcap SIGNAL \ 2$

This code is used in chunk 22.

Uses SIGNAL 20b.

65d *⟨Nest top-level root lines as Z nodes 65d⟩*≡
 $_ \leftarrow (gz \ 1\phi _)' (t[i]=Z) < i \leftarrow \underline{1} d=0$
 $'Non-Z top-level node' assert \ t[\underline{1} p = i \neq p] = Z:$

This code is used in chunk 22.

66a *⟨Parse :Namespace syntax 66a⟩*≡
 nss←nε<':NAMESPACE' ♦ nse←nε<':ENDNAMESPACE'
 ERR←':NAMESPACE KEYWORD MAY ONLY APPEAR AT BEGINNING OF A LINE'
 Zv.≠t≠1φnss:ERR □SIGNAL 2
 ERR←'NAMESPACE DECLARATION MAY HAVE ONLY A NAME OR BE EMPTY'
 v≠(Z≠t≠1φnss)^(V≠t≠1φnss)∨Z≠t≠2φnss:ERR □SIGNAL 2
 ERR←':ENDNAMESPACE KEYWORD MUST APPEAR ALONE ON A LINE'
 v≠Z≠t≠1 -1v.φcnse:ERR □SIGNAL 2
 t[nsi←1φnss]←M ♦ t[nei←1φnse]←-M
 n[i]←n[1+i←1(t=M)∧V=1φt] ♦ end[nsi]←end[nei]
 x←1p=1≠p ♦ d←+λ(t[x]=M)+-t[x]=-M
 0≠φd:':NAMESPACE KEYWORD MISSING :ENDNAMESPACE PAIR'□SIGNAL 2
 p[x]←x[D2P -1φd]

 A Delete unnecessary namespace nodes from the tree, leave only M's
 msk←~nssv((-1φnss)∧t=V)∨nsev1φnse
 t k n pos end≠←msk ♦ p←(1~msk)(1-1+1)msk≠p
 This code is used in chunk 22.
 Uses SIGNAL 20b.

In the parser, the *xn* and *xt* fields are not part of the AST proper, but form an auxiliary analysis that is exceptionally useful, and so we include this as a part of the output of the parser. After parsing a module, we want to extract out the top-level bindings and what their types are, which we can then use to feed into things like the linker and other areas that might need to know what names are available in a given module. Top-level bindings are identified as bindings that appear as a part of an initialization function, also known as F0.

66b *⟨Compute parser exports 66b⟩*≡
 msk←(t=B)∧k[I@{t[ω]≠F}≠p]=0
 xn←(Op<''),msk≠n ♦ xt←msk≠k
 This code is used in chunk 17.
 Defines:
 xn, used in chunk 20a.
 xt, used in chunk 20a.

66c *⟨Record exported top-level bindings 66c⟩*≡
 xi←1(t=B)∧k[r]=0
 This code is used in chunk 23.
 Defines:
 xi, used in chunks 23–25.

66d *⟨Node ↔ Generator mapping 24c⟩*+≡
 gck,←<F 0
 gcν,←<'Fz'
 This code is used in chunk 25b.

67 $\langle \text{Node-specific code generators } 24e \rangle + \equiv$

```

Fz ← {id ← 5; α ← awc ← v ≠ (3[]x) { (ω ∈ A 0) ∨ (ω = E) ∧ α > 0 } 2[]x ← Q; ≠ ω
z ← c 'int init', id, ' = 0;'
z, ← c '
z, ← c 'EXPORT int'
z, ← c 'init(void)'
z, ← c '{'
z, ← c ' return fn', id, '(NULL, NULL, NULL, NULL);'
z, ← c '}'
z, ← c '
z, ← c 'int'
z, ← c 'fn', id, '(struct array **z, '
z, ← c '    struct array *l, struct array *r, void *fv[])'
z, ← c '{'
z, ← c '    void    *stk[128];'
z, ← c '    void    **stkhd;'
z, ← c '    void    *a, *w;'
z, ← c '    struct closure *c;'
z, ← c '
z, ← c '    if (init', id, ') '
z, ← c '        return 0;'
z, ← c '
z, ← c '    stkhd = &stk[0];'
z, ← c '    init', id, ' = 1;'
z, ← c '    cdf_init();'
z, ← c '
z, ← c '    , " , ≠ dis " ω
z, ← c '        return 0;'
z, ← c '}'
z, ← c '
z}

```

This code is used in chunk 25b.

68a $\langle \textit{init.c}$ 68a) \equiv

```
#include "codfns.h"

int
init(void);

EXPORT int
cdf_init(void)
{
  return init();
}
```

Root chunk (not used in this document).

Defines:

`init.c`, used in chunk 68b.

Uses `codfns` 7, `codfns.h` 33, and `EXPORT` 34b.

68b $\langle \textit{Tangle Commands}$ 8) \equiv

```
echo "Tangling rtm/init.c..."
notangle -R'init.c' codfns.nw > rtm/init.c
```

This code is used in chunk 83.

Uses `codfns` 7 and `init.c` 68a.

6.16.3 Structured Programming Statements

68c $\langle \textit{Verify that all structured statements appear within trad-fns}$ 68c) \equiv

```
msk←kwseKW~':NAMESPACE' ':ENDNAMESPACE' ':SECTION' ':ENDSECTION'
v≠msk←msk^~km≠tm:{
  msg←2'STRUCTURED STATEMENTS MUST APPEAR WITHIN TRAD-FNS'
  msg SIGNAL ε{x+end[ω]-x←pos[ω]}''_1km\msk
}θ
```

This code is used in chunk 22.

Uses `SIGNAL` 20b.

68d $\langle \textit{Convert M nodes to FO nodes}$ 68d) \equiv

```
t←F@{t=M}t
```

This code is used in chunk 22.

7 Runtime Primitives

7.1 Addition/Identity

68e $\langle \textit{Symbol} \leftrightarrow \textit{Name mapping}$ 24b) \equiv

```
syms,←c, '+' ◇ nams,←c'add'
```

This code is used in chunk 25b.

7.2 And (Logical)

69a $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \wedge ' \diamond \text{nams}, \leftarrow c, \text{'and'}$

This code is used in chunk 25b.

7.3 Bracket

69b $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' [' \diamond \text{nams}, \leftarrow c, \text{'brk'}$

This code is used in chunk 25b.

7.4 Catenate (First/Last Axis)

69c $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ', ' \diamond \text{nams}, \leftarrow c, \text{'cat'}$
 $\text{syms}, \leftarrow c, ' ; ' \diamond \text{nams}, \leftarrow c, \text{'ctf'}$

This code is used in chunk 25b.

7.5 Circle/Trigonometrics

69d $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' o ' \diamond \text{nams}, \leftarrow c, \text{'cir'}$

This code is used in chunk 25b.

7.6 Commute

69e $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \rightleftarrows ' \diamond \text{nams}, \leftarrow c, \text{'com'}$

This code is used in chunk 25b.

7.7 Compose

69f $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \circ ' \diamond \text{nams}, \leftarrow c, \text{'jot'}$

This code is used in chunk 25b.

7.8 Convolve

69g $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \square \text{CONV} ' \diamond \text{nams}, \leftarrow c, \text{'conv'}$

This code is used in chunk 25b.

7.9 Decode

70a $\langle \textit{Symbol} \leftrightarrow \textit{Name mapping} \ 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \bot ' \ \diamond \ \text{nams}, \leftarrow c ' \text{dec} '$
 This code is used in chunk 25b.

7.10 Disclose

70b $\langle \textit{Symbol} \leftrightarrow \textit{Name mapping} \ 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \supset ' \ \diamond \ \text{nams}, \leftarrow c ' \text{dis} '$
 This code is used in chunk 25b.

7.11 Division/Reciprocal

70c $\langle \textit{Symbol} \leftrightarrow \textit{Name mapping} \ 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \div ' \ \diamond \ \text{nams}, \leftarrow c ' \text{div} '$
 This code is used in chunk 25b.

7.12 Drop

70d $\langle \textit{Symbol} \leftrightarrow \textit{Name mapping} \ 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \downarrow ' \ \diamond \ \text{nams}, \leftarrow c ' \text{drp} '$
 This code is used in chunk 25b.

7.13 Each

70e $\langle \textit{Symbol} \leftrightarrow \textit{Name mapping} \ 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \cdot ' \ \diamond \ \text{nams}, \leftarrow c ' \text{map} '$
 This code is used in chunk 25b.

7.14 Enclose

70f $\langle \textit{Symbol} \leftrightarrow \textit{Name mapping} \ 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' c ' \ \diamond \ \text{nams}, \leftarrow c ' \text{par} '$
 This code is used in chunk 25b.

7.15 Encode

70g $\langle \textit{Symbol} \leftrightarrow \textit{Name mapping} \ 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \top ' \ \diamond \ \text{nams}, \leftarrow c ' \text{enc} '$
 This code is used in chunk 25b.

7.16 Equal

71a $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, '=' \diamond \text{nams}, \leftarrow c 'eql'$

This code is used in chunk 25b.

7.17 Exponent

71b $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, '*' \diamond \text{nams}, \leftarrow c 'exp'$

This code is used in chunk 25b.

7.18 Factorial/Binomial

71c $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, '!' \diamond \text{nams}, \leftarrow c 'fac'$

This code is used in chunk 25b.

7.19 Fast Fourier Transforms

71d $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, '\square FFT' \diamond \text{nams}, \leftarrow c 'fft'$

This code is used in chunk 25b.

71e $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, '\square IFFT' \diamond \text{nams}, \leftarrow c 'ift'$

This code is used in chunk 25b.

7.20 Find

71f $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, '\underline{_}' \diamond \text{nams}, \leftarrow c 'fnd'$

This code is used in chunk 25b.

7.21 Grade Down

71g $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, '\Psi' \diamond \text{nams}, \leftarrow c 'gdd'$

This code is used in chunk 25b.

7.22 Grade Up

72a $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ←c, 'Δ' ◇ nams, ←c 'gdu'`

This code is used in chunk 25b.

7.23 Greater Than

72b $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ←c, '>' ◇ nams, ←c 'gth'`

This code is used in chunk 25b.

7.24 Greater Than or Equal

72c $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ←c, '≥' ◇ nams, ←c 'gte'`

This code is used in chunk 25b.

7.25 Index

72d $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ←c, '[]' ◇ nams, ←c 'sqd'`

This code is used in chunk 25b.

7.26 Index Generator

72e $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ←c, 'ι' ◇ nams, ←c 'iot'`

This code is used in chunk 25b.

7.27 Inner Product

72f $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ←c, '.' ◇ nams, ←c 'dot'`

This code is used in chunk 25b.

7.28 Intersection

72g $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ←c, 'n' ◇ nams, ←c 'int'`

This code is used in chunk 25b.

7.29 Left

73a $\langle \textit{Symbol} \leftrightarrow \textit{Name mapping} \ 24b \rangle + \equiv$
 $\text{syms}, \leftarrow, ' \rightarrow ' \diamond \text{nams}, \leftarrow, ' \textit{left} '$

This code is used in chunk 25b.

7.30 Less Than

73b $\langle \textit{Symbol} \leftrightarrow \textit{Name mapping} \ 24b \rangle + \equiv$
 $\text{syms}, \leftarrow, ' < ' \diamond \text{nams}, \leftarrow, ' \textit{lt} h '$

This code is used in chunk 25b.

7.31 Less Than or Equal

73c $\langle \textit{Symbol} \leftrightarrow \textit{Name mapping} \ 24b \rangle + \equiv$
 $\text{syms}, \leftarrow, ' \leq ' \diamond \text{nams}, \leftarrow, ' \textit{lte} '$

This code is used in chunk 25b.

7.32 Logarithm

73d $\langle \textit{Symbol} \leftrightarrow \textit{Name mapping} \ 24b \rangle + \equiv$
 $\text{syms}, \leftarrow, ' \otimes ' \diamond \text{nams}, \leftarrow, ' \textit{log} '$

This code is used in chunk 25b.

7.33 Match

73e $\langle \textit{Symbol} \leftrightarrow \textit{Name mapping} \ 24b \rangle + \equiv$
 $\text{syms}, \leftarrow, ' \equiv ' \diamond \text{nams}, \leftarrow, ' \textit{eqv} '$

This code is used in chunk 25b.

7.34 Matrix Division

73f $\langle \textit{Symbol} \leftrightarrow \textit{Name mapping} \ 24b \rangle + \equiv$
 $\text{syms}, \leftarrow, ' \boxdiv ' \diamond \text{nams}, \leftarrow, ' \textit{mdv} '$

This code is used in chunk 25b.

7.35 Maximum/Ceiling

73g $\langle \textit{Symbol} \leftrightarrow \textit{Name mapping} \ 24b \rangle + \equiv$
 $\text{syms}, \leftarrow, ' \lceil ' \diamond \text{nams}, \leftarrow, ' \textit{max} '$

This code is used in chunk 25b.

7.36 Membership

74a $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \epsilon ' \diamond \text{nams}, \leftarrow c, ' \text{mem} '$

This code is used in chunk 25b.

7.37 Minimum/Floor

74b $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \lfloor ' \diamond \text{nams}, \leftarrow c, ' \text{min} '$

This code is used in chunk 25b.

7.38 Multiplication

74c $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \times ' \diamond \text{nams}, \leftarrow c, ' \text{mul} '$

This code is used in chunk 25b.

7.39 Nest/Partition

74d $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \subseteq ' \diamond \text{nams}, \leftarrow c, ' \text{nst} '$

This code is used in chunk 25b.

7.40 Not

74e $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \sim ' \diamond \text{nams}, \leftarrow c, ' \text{not} '$

This code is used in chunk 25b.

7.41 Not And (Logical)

74f $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \tilde{\wedge} ' \diamond \text{nams}, \leftarrow c, ' \text{nan} '$

This code is used in chunk 25b.

7.42 Not Equal

74g $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \neq ' \diamond \text{nams}, \leftarrow c, ' \text{neq} '$

This code is used in chunk 25b.

7.43 Not Match

75a $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \neq ' \diamond \text{nams}, \leftarrow c ' nqv '$

This code is used in chunk 25b.

7.44 Not Or (Logical)

75b $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \tilde{v} ' \diamond \text{nams}, \leftarrow c ' nor '$

This code is used in chunk 25b.

7.45 Or (Logical)

75c $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' v ' \diamond \text{nams}, \leftarrow c ' lor '$

This code is used in chunk 25b.

7.46 Outer Product

75d $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \circ . ' \diamond \text{nams}, \leftarrow c ' oup '$

This code is used in chunk 25b.

7.47 Power

75e $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \times ' \diamond \text{nams}, \leftarrow c ' pow '$

This code is used in chunk 25b.

7.48 Rank

75f $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' \ddot{o} ' \diamond \text{nams}, \leftarrow c ' rnk '$

This code is used in chunk 25b.

7.49 Reduce

75g $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
 $\text{syms}, \leftarrow c, ' / ' \diamond \text{nams}, \leftarrow c ' red '$

This code is used in chunk 25b.

76a $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ← c, 'f' ⋄ nams, ← c 'rdf'`
 This code is used in chunk 25b.

7.50 Roll

76b $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ← c, '?' ⋄ nams, ← c 'rol'`
 This code is used in chunk 25b.

7.51 Rotate (First/Last Axis)

76c $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ← c, 'φ' ⋄ nams, ← c 'rot'`
`syms, ← c, 'θ' ⋄ nams, ← c 'rtf'`
 This code is used in chunk 25b.

7.52 Residue

76d $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ← c, '|' ⋄ nams, ← c 'res'`
 This code is used in chunk 25b.

7.53 Right

76e $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ← c, 'r' ⋄ nams, ← c 'rgt'`
 This code is used in chunk 25b.

76f $\langle \text{APL Primitives } 76f \rangle \equiv$
`rgt ← {ω}`
 This code is used in chunk 32a.
 Defines:
 rgt, used in chunk 80.

7.54 Scalar Each

76g $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ← c, '%s' ⋄ nams, ← c 'scl'`
 This code is used in chunk 25b.

7.55 Scan

77a $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ←c, '\ ' ◊ nams, ←c 'scn'`
 This code is used in chunk 25b.

77b $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ←c, 'x' ◊ nams, ←c 'scf'`
 This code is used in chunk 25b.

7.56 Shape

77c $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ←c, 'ρ' ◊ nams, ←c 'rho'`
 This code is used in chunk 25b.

7.57 Subtraction

77d $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ←c, '-' ◊ nams, ←c 'sub'`
 This code is used in chunk 25b.

7.58 Take

77e $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ←c, 't' ◊ nams, ←c 'tke'`
 This code is used in chunk 25b.

7.59 Transpose

77f $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ←c, 'Q' ◊ nams, ←c 'trn'`
 This code is used in chunk 25b.

7.60 Union

77g $\langle \text{Symbol} \leftrightarrow \text{Name mapping } 24b \rangle + \equiv$
`syms, ←c, 'u' ◊ nams, ←c 'unq'`
 This code is used in chunk 25b.

8 Utilities

8.1 Must haves

There are some APL functions that are so critical as to be worthy of primitive status.

- Indexing
- Under
- Assert

```
78  ⟨Must Have APL Utilities 78⟩≡
    I←{(cω)[]α}
    U←{α←⊥ ⋄ ωω×¯1⊥α ααöωω ω}
    assert←{
    α←'assertion failure'
    0∈ω:⊥'α [] SIGNAL 8'
    1:shy←0
    }
```

This code is used in chunk 7.

Defines:

 assert, used in chunk 22.

Uses SIGNAL 20b.

8.2 AST Pretty-printing

```

79  ⟨Pretty-printing AST trees 79⟩≡
    dct←{α[(2×2≠/n,0)+(1↑≠m)+m+n←φv\φm←' '≠αα ω]ωω ω}
    dlk←{((x□ρω)↑[x←2|1+ωω]α),[ωω]αα@(c0 0)×('┐'⇒ω)┐ω}

    dwh←{
      z←↗/((≠''α),''c┐/≠○φ''α)↑''α
      ω('┐'dlk 1)' |┐┐┐'(0□φ)dct,z
    }
    dwv←{
      z←{α,' ',ω}/(1+┐/≠''α){α↑ω↗''| '↑≠φω}''α
      ω('┐'dlk 0)' ┐┐┐| '(0□┐)dct(↗1┐┐)z
    }

    lb3←{
      α←ι≠ω
      z←(NΔ{α[ω]}@2┐(2>ω){α[|ω]}@{0>ω}@4↑>ω)[α; ]
      '( ',''' )',''≠{α,','; ',ω}≠≠z
    }

    pp3←{
      α←'o' ◇ lbl←αρ≠ω
      d←(ι≠ω)≠ω ◇ _←{z┐d+←ω≠z←α[ω]}×≡≠ω
      lyr←{
        i←┐α=d
        k v←┐φωω[i ],○c┐i
        (ω○{α[ω]}''v)αα''@k┐ω
      }ω
      (ω=ι≠ω)≠αα lyr≠(1+ι┐/d),cφ○↗○≠''lbl
    }

```

This code is used in chunk 7.

Defines:

dct, never used.
 dlk, never used.
 dwh, never used.
 dwv, never used.
 lb3, never used.
 pp3, never used.

8.3 Debugging utilities

The following utilities help to improve quality of life when working with the Co-dfns source code.

The `DISPLAY` function is taken from <https://dfns.dyalog.com> and helps to make debugging easier by allowing us to thread `DISPLAY` calls into expressions. I prefer to do something like this:

```
... {ω←⊠+#.DISPLAY ω} ...
```

The function itself returns the character rendering of the code, so the above little expression is one that I use to insert and do debugging within an expression.

```
80 (DISPLAY Utility 80)≡
  DISPLAY←{
    ⊠IO ⊠ML←0
    α←1 ⋄ chars←α>'..'''|-' ' ⊠ ⊠ |-'
    tl tr bl br vt hz←chars
    box←{
      vrt hrz←(⊖1+ρω)ρ⊖vt hz
      top←(hz,'⊖')[⊖1↑α],hrz
      bot←(α),hrz
      rgt←tr,vt,vrt,br
      lax←(vt,'⊕')[⊖1↑1↑α],⊖c vrt
      lft←⊕tl,(↑lax),bl
      lft,(top;ω;bot),rgt
    }
    deco←{α←type open ω ⋄ α,axes ω}
    axes←{(-2⌈ρρω)↑1+×ρω}
    open←{(1⌈ρρω)ρω}
    trim←{(~1 1⊆^ω≠' ')/ω}
    type←{{(1=ρω)⊃'+ω}⊃,char''ω}
    char←{⊖≡ρω:hz ⋄ (ω∈'-' ,⊠D)⊃'#~'}⊃⊥
    line←{(6≠10|⊠DR' 'ω)⊃' -'}
    {
      0≡ω:' ' ;(open ⊠FMT ω);line ω
      1 ⊖≡(≡ω)(ρω):'∇' 0 0 box ⊠FMT ω
      1≡ω:(deco ω)box open ⊠FMT open ω
      ('ε'deco ω)box trim ⊠FMT ∇''open ω
    }ω
  }
```

Root chunk (not used in this document).

Defines:

`DISPLAY`, used in chunk 81.

Uses `box` 56d, `rgt` 76f, `⊠IO` 10a, and `⊠ML` 10a.

I also define a function `PP` that encapsulates the above usage pattern that I like to use, making the whole thing less verbose and a little more convenient.

81a `<PP Utility 81a>≡`
`PP←{ω←⍋#.DISPLAY ω}`
 Root chunk (not used in this document).
 Defines:
`PP`, used in chunks 27 and 81b.
 Uses `DISPLAY 80`.

Both of these function exist outside of the `codfns` namespace and so they get their own files inside of the `src\` directory.

81b `<Tangle Commands 8>+≡`
`echo "Tangling src/DISPLAY.aplf..."`
`notangle -R'[[DISPLAY]] Utility' codfns.nw > src/DISPLAY.aplf`

`echo "Tangling src/PP.aplf..."`
`notangle -R'[[PP]] Utility' codfns.nw > src/PP.aplf`
 This code is used in chunk 83.
 Defines:
`DIRECTORY`, never used.
`PP.aplf`, never used.
 Uses `codfns 7`, `DISPLAY 80`, `PP 81a`, and `src 88`.

8.4 Reading and Writing Files

It is helpful to be able to easily write files to disk, and the following `put` and `tie` utilities help us to do so when we want to. These are pretty standard, but they could maybe be replaced by `⍋INPUT` or something like that.

81c `<Basic tie and put utilities 81c>≡`
`tie←{`
`0::⍋SIGNAL ⍋EN`
`22::ω ⍋NCREATE 0`
`0 ⍋NRESIZE ω ⍋NTIE 0`
`}`

`put←{`
`s←(¯128+256|128+'UTF-8'⍋UCS ω)⍋NAPPEND(t←tie α)83`
`1:r←s⍋NUNTIE t`
`}`

This code is used in chunks 7 and 87b.
 Defines:
`put`, used in chunks 26, 87b, and 88.
`tie`, used in chunk 87b.
 Uses `SIGNAL 20b`.

8.5 XML Rendering

82a $\langle XML\ Rendering\ 82a \rangle \equiv$

```

  Xml ← {α ← 0
    ast ← α {d i ← P2D ⊃ ω ⋄ i ∘ {ω[α]}''(c d), 1 ↓ α ↓ ω} * (0 ≠ α) ⊢ ω
    d t k n ← 4 ↑ ast
    cls ← NΔ[t], ''(' - . . '[1 + × k]), ''⌘'' | k
    fld ← {((≠ ω) ↑ 3 ↓ f Δ), ⌘ ω}'' ↓ ⌘ ↑ 3 ↓ ast
    □ XML ⌘ ↑ d cls (c ' ') fld
  }

```

This code is used in chunk 7.

Defines:

Xml, never used.

8.6 Detecting the Operating System

It is quite helpful to be able to easily detect the operating system that we are on. This turns out to be helpful in more areas than just the compiler.

82b $\langle The\ opsys\ utility\ 82b \rangle \equiv$

```

  opsys ← {ω ⊃ ⌘ 'Win' 'Lin' 'Mac' ⌘ c 3 ↑ ⊃ ' . ' □ WG 'APLVersion'}

```

This code is used in chunks 7, 84c, and 86d.

Defines:

opsys, used in chunks 26, 84c, and 86d.

9 Developer Infrastructure

9.1 Building the Compiler

The Co-dfns compiler is written, developed, and distributed as a literate program. For more information about literate programming, see the resources available at <http://literateprogramming.com/>. We use noweb as our preferred literate programming tool because it is eminently simple, while still handling the majority of our needs and producing high quality output in L^AT_EX format with all the important elements of literate programming, including live hyperlinking and cross-references.

9.1.1 Tangling the Source

The process of tangling produces the executable source code for the compiler. Importantly, the tangled output is *not* meant to be used as the primary means of reading or debugging the source. Instead, it is meant primarily as the machine readable version of the code only.

With noweb, we need to invoke `notangle` once for each of the chunks that we wish to use to produce an output file. To make this easy, we build up a script to do this work for us.

For Linux and Mac, the following bash script creates these files. We use a separate chunk that we build up incrementally throughout the rest of this document as a record of all the chunks that we should create. Notice that we explicitly tangle the `TANGLE.sh` file as the last thing that we do; this helps to ensure that we are reliably executing the rest of the script before changing the contents of the file, as some systems will be affected and change execution behavior in strange ways if we change the `TANGLE.sh` file early on in the execution of the file.

```
83 <TANGLE.sh 83>≡
    #!/bin/bash

    <Tangle Commands 8>

    echo "Tangling TANGLE.sh..."
    notangle -R'[[TANGLE.sh]]' codfns.nw > TANGLE.sh
Root chunk (not used in this document).
Defines:
    TANGLE.sh, used in chunk 84a.
Uses codfns 7 and TANGLE 84c.
```

On Windows, the best way that we have found to do this is by installing noweb using the Cygwin project and then calling `TANGLE.sh` from a local `TANGLE.bat` file. This document assumes that you have already successfully built and installed via Cygwin a working Icon-driven noweb installation.

Users who prefer to work in a UNIX fashion via Cygwin or some other subsystem on Windows can follow the build scripts directly. For developers who prefer to work in a primarily Windows environment, the following `TANGLE.bat` build script assists in handling the calls into Cygwin so that you do not need to have a Cygwin terminal open all the time.

84a `<TANGLE.bat 84a>≡`
`set SH=C:\cygwin64\bin\bash.exe -l -c`
`%SH% "cd $OLDPWD && ./TANGLE.sh"`

Root chunk (not used in this document).

Defines:

`TANGLE.bat`, used in chunk 84b.

Uses `TANGLE 84c` and `TANGLE.sh 83`.

84b `<Tangle Commands 8>+≡`
`echo "Tangling TANGLE.bat..."`
`notangle -R'[[TANGLE.bat]]' codfns.nw > TANGLE.bat`

This code is used in chunk 83.

Uses `codfns 7`, `TANGLE 84c`, and `TANGLE.bat 84a`.

When tangled to the `TANGLE.aplf` file, the following script enables the user to simply type `TANGLE` within a Dyalog APL session to update the code tree from within Dyalog itself. This is much more convenient than keeping a Cygwin Terminal session open along with a Dyalog APL session while programming.

Note: this command expects to be run from within the root of the repository, not from, say, within the testing directory.

84c `<TANGLE 84c>≡`
`TANGLE;opsys`
`<The opsys utility 82b>`
`□CMD opsys '.\TANGLE.bat' './TANGLE.sh' './TANGLE.sh'`

Root chunk (not used in this document).

Defines:

`TANGLE`, used in chunks 83 and 84.

Uses `opsys 82b`.

84d `<Tangle Commands 8>+≡`
`echo "Tangling TANGLE.aplf..."`
`notangle -R'[[TANGLE]]' codfns.nw > src/TANGLE.aplf`

This code is used in chunk 83.

Defines:

`TANGLE.aplf`, never used.

Uses `codfns 7`, `src 88`, and `TANGLE 84c`.

9.1.2 Weaving the Source

Weaving is the process by which we produce the final printed output of this document, intended for reading and general human consumption. We rely on the \LaTeX typesetting system to do this. Moreover, because we make heavy use of UTF-8 and prefer to have our own fonts installed and used, it is necessary to use the `xelatex` system instead of the typical \LaTeX engine. In order to get the indexing right, we must run the engine twice. The first run will update the indexing files that will be picked up on the second run and incorporated into the final document. Note, we have tried to use the `lua-latex` engine, which in theory should work just as well as the `xelatex` engine, but we get a strange error relating to noweb's style file, so we stick with `xelatex` for now.

Running this script also depends on having the appropriate fonts installed. In this case, please ensure that the following fonts are installed in your Windows font system so that they can be picked up by the \TeX engine.

- Libre Baskerville (Regular, Italic, Bold)
- APL385 Unicode
- Lucida Sans Unicode
- Cambria Math

If you do not wish to use these fonts, edit the font specifications at the top of `codfns.nw` to the fonts that you do wish to use.

Note the use of `-delay -index` for options. We want to generate indexing, but we also need to make sure that we can use some of our own packages in the system,

Note: this command expects to be run from within the root of the repository, not from, say, within the testing directory.

```
85 <WEAVE.sh 85>≡
    #!/bin/bash
    mkdir -p woven
    noweave -delay -index codfns.nw > woven/codfns.tex
    cd woven
    xelatex --shell-escape codfns
    xelatex --shell-escape codfns
```

Root chunk (not used in this document).

Defines:

`WEAVE.sh`, used in chunk 86.

Uses `codfns` 7.

```
86a  <Tangle Commands 8>+≡
      echo "Tangling WEAVE.sh..."
      notangle -R'[[WEAVE.sh]]' codfns.nw > WEAVE.sh
```

This code is used in chunk 83.

Uses codfns 7, WEAVE 86d, and WEAVE.sh 85.

And just like the tangling code, we want to define a TANGLE.bat batch file to call the Cygwin environment from Windows.

```
86b  <WEAVE.bat 86b>≡
      set SH=C:\cygwin64\bin\bash.exe -l -c
      %SH% "cd $OLDPWD && ./WEAVE.sh"
```

Root chunk (not used in this document).

Defines:

WEAVE.bat, used in chunk 86c.

Uses WEAVE 86d and WEAVE.sh 85.

```
86c  <Tangle Commands 8>+≡
      echo "Tangling WEAVE.bat..."
      notangle -R'[[WEAVE.bat]]' codfns.nw > WEAVE.bat
```

This code is used in chunk 83.

Uses codfns 7, WEAVE 86d, and WEAVE.bat 86b.

Like the *<TANGLE Command (never defined)>*, the following command, when tangled to the WEAVE.aplf file enables weaving in a the Dyalog APL session by executing the WEAVE command.

```
86d  <WEAVE 86d>≡
      WEAVE;opsys
      <The opsys utility 82b>
      □CMD opsys '.\WEAVE.bat' './WEAVE.sh' './WEAVE.sh'
```

Root chunk (not used in this document).

Defines:

WEAVE, used in chunk 86.

Uses opsys 82b.

```
86e  <Tangle Commands 8>+≡
      echo "Tangling src/WEAVE.aplf..."
      notangle -R'[[WEAVE]]' codfns.nw > src/WEAVE.aplf
```

This code is used in chunk 83.

Defines:

WEAVE.aplf, never used.

Uses codfns 7, src 88, and WEAVE 86d.

9.2 Building the Runtime

One of our goals with the Co-dfns runtime is to write as much of it as possible in APL. This means that we want to have at minimum a very small kernel that has been written in C, while most of the rest of the code is implemented in some APL files. This leads to a three part breakdown of the process to build the runtime.

87a *⟨Build the runtime 87a⟩*≡
 ⟨Compile the primitives in prim.apln 88⟩
 ⟨Build codfns.dll DLL 89a⟩
 ⟨Copy the runtime files into tests\ 89b⟩

This code is used in chunk 87b.

We define the command `MKΔRTM` to build the runtime. This command takes a path to the root directory of the Co-dfns repository; this is to allow us to rebuild the runtime from anywhere in the system if we so choose.

87b *⟨MKΔRTM 87b⟩*≡
 `MKΔRTM path;put;tie;src;vsbat;vsc;wsd`

⟨Basic tie and put utilities 81c⟩
⟨Build the runtime 87a⟩

Root chunk (not used in this document).

Defines:

`MKΔRTM`, used in chunk 87c.

Uses `put 81c`, `src 88`, `tie 81c`, `vsbat 89a`, `vsc 89a`, and `wsd 89a`.

This file is another of our external utilities that exists outside of the `codfns` namespace, so it gets its own file in `src\`.

87c *⟨Tangle Commands 8⟩*+≡
 `echo "Tangling src/MKΔRTM.aplf..."`
 `notangle -R'[[MKΔRTM]]' codfns.nw > src/MKΔRTM.aplf`

This code is used in chunk 83.

Defines:

`MKΔRTM.aplf`, never used.

Uses `codfns 7`, `MKΔRTM 87b`, and `src 88`.

The first step we must take is producing an appropriate C file that contains the primitives that we have defined in `prim.apln`. This means that we want to only compile the code in `prim.apln` as far as producing the C code. Since we do not have a full blown runtime yet, we will be compiling the `prim.c` file along with the rest of the runtime code, instead of the normal build process, which assumes that we already have a working runtime. This means that we only invoke the GC TT PS passes of the compiler pipeline, while avoiding the CC pass. We use the SALT system to load the source from `prim.apln` and then run the compiler passes that we want before storing the resulting code in the `rtm\prim.c` file.

```
88  <Compile the primitives in prim.apln 88>≡
      src←SRC SE.SALT.Load path,'\rtm\prim.apln'
      (path,'\rtm\prim.c')put codfns.{GC TT PS ω}src
```

This code is used in chunk 87a.

Defines:

`src`, used in chunks 8, 13, 16b, 23, 81b, 84d, 86, and 87.

Uses `codfns` 7, `prim` 32a, PS 17, and `put` 81c.

Once we have the `rtm\prim.c` file written appropriately, we can run the main compiler process. For simplicity, we just compile all of the `.c` files that are found in the `rtm\` subdirectory. We must ensure that we are appropriately invoking our ArrayFire dependencies as well as producing the appropriate debugging symbols most of the time.

```
89a <Build codfns.dll DLL 89a>≡
    vsbat←#.codfns.VSΔPATH
    vsbat, '\VC\Auxiliary\Build\vcvarsall.bat'
    wsd←path, '\

    vsc←'%comspec% /C "', vsbat, '" amd64'
    vsc,←' && cd "', wsd, '\rtm"'
    vsc,←' && cl /MP /W3 /wd4102 /wd4275'
    vsc,←' /Od /Zc:inline /Zi /FS'
    vsc,←' /Fo".\\" /Fd"codfns.pdb"'
    vsc,←' /WX /MD /EHsc /nologo'
    vsc,←' /I"%AF_PATH%\include"'
    vsc,←' /D"NOMINMAX" /D"AF_DEBUG" /D"EXPORTING"'
    vsc,←' "*.c" /link /DLL /OPT:REF'
    vsc,←' /INCREMENTAL:NO /SUBSYSTEM:WINDOWS'
    vsc,←' /LIBPATH:"%AF_PATH%\lib"'
    vsc,←' /DYNAMICBASE "af",codfns.AFΔLIB, '.lib"'
    vsc,←' /OPT:ICF /ERRORREPORT:PROMPT'
    vsc,←' /TLBID:1 /OUT:"codfns.dll"'
```

This code is used in chunk 87a.

Defines:

`vsbat`, used in chunks 26 and 87b.

`vsc`, used in chunks 26, 87b, and 89b.

`wsd`, used in chunks 87b and 89b.

Uses `AFΔLIB` 11, `codfns` 7, `EXPORTING` 34b, and `VSΔPATH` 12.

Finally, in order to write up the test harness to work right, we must copy the appropriate runtime files into the `tests\` directory so that we can find them when we finally start running our code there.

```
89b <Copy the runtime files into tests\ 89b>≡
    □CMD □←vsc
    □CMD □←'copy "', wsd, 'rtm\codfns.h" "', wsd, 'tests\'
    □CMD □←'copy "', wsd, 'rtm\codfns.exp" "', wsd, 'tests\'
    □CMD □←'copy "', wsd, 'rtm\codfns.lib" "', wsd, 'tests\'
    □CMD □←'copy "', wsd, 'rtm\codfns.pdb" "', wsd, 'tests\'
    □CMD □←'copy "', wsd, 'rtm\codfns.dll" "', wsd, 'tests\'
```

This code is used in chunk 87a.

Uses `codfns` 7, `codfns.h` 33, `vsc` 89a, and `wsd` 89a.

9.3 Loading the Compiler

In order to load the compiler into an APL session as well as all the development utilities, we assume that you have first managed to either load up a session with a bootstrapped version of the `TANGLE` command or that you already have a tangled `src\` directory. If the `src\` directory has not yet been created by running the `TANGLE` command, then this must be done before loading the compiler system. After tangling, the compiler can be loaded using the provided `LOAD` shortcut. This shortcut is meant to use the Dyalog Link system for hot-loading the files in `src\` into the root namespace. We do so through the following link command:

```
Link.Create # src -source=dir -watch=dir
```

This means that we want to link the `src\` directory into the `#` namespace, but we also want to make sure that we only pull changes that come from the filesystem. This is because we are editing the code via the `WEB` document, and we do not want to risk having some intermediate representation that isn't accurate and that doesn't flow the right way; we want all appropriate changes to begin in the `WEB` document and then, and only then, flow into the session. This also allows us to make some modifications to the code for testing and experimentation inside of the session without consideration for the code outside of the session, and such changes will be removed or forgotten on the next `TANGLE` command.

To set this up, we also ensure that we begin our work within the root Co-dfns repository directory, as this is where we expect to run the `TANGLE` and `WEAVE` commands.

There is unfortunately only a limited range of possibilities for linking in a new directory as we wish to do. The method we choose to use is launching a fresh Dyalog APL session and then using an `LX` expression from the command line to do the actual linking using the `⎕SE.UCMD` functionality. I personally find this to be rather hackish, and I hope that an alternative approach to doing this will show up in the near future. Nonetheless, the arguments that we pass to `dyalog.exe` look something like this:

```
LX="⎕SE.UCMD'Link.Create # src -source=dir -watch=dir'"
```

If you do not use the `LOAD` shortcut, you can use the above command to do the linking manually.

10 Index

10.1 Chunks

⟨* 7⟩
 ⟨DISPLAY *Utility* 80⟩
 ⟨MKΔRTM 87b⟩
 ⟨PP *Utility* 81a⟩
 ⟨TANGLE.bat 84a⟩
 ⟨TANGLE.sh 83⟩
 ⟨TANGLE 84c⟩
 ⟨TEST 16a⟩
 ⟨WEAVE.bat 86b⟩
 ⟨WEAVE.sh 85⟩
 ⟨WEAVE 86d⟩
 ⟨*Adjust AST for output* 18b⟩
 ⟨*Anchor variables to earliest binding in the matching frame* 61f⟩
 ⟨APL *Primitives* 76f⟩
 ⟨*AST Record Structure* 15b⟩
 ⟨*Basic tie and put utilities* 81c⟩
 ⟨*Box definitions* 57a⟩
 ⟨*box.c* 58a⟩
 ⟨*Build codfns.dll DLL* 89a⟩
 ⟨*Build the runtime* 87a⟩
 ⟨*C runtime declarations* 38a⟩
 ⟨*C runtime enumerations* 36b⟩
 ⟨*C runtime includes* (never defined)⟩
 ⟨*C runtime macros* 34b⟩
 ⟨*C runtime structures* 36a⟩
 ⟨*Cell definitions* 37⟩
 ⟨*Cell release cases* 39c⟩
 ⟨*Cell type names* 36c⟩
 ⟨*cell.c* 41a⟩
 ⟨*Check and mask the strings* 49a⟩
 ⟨*Check for out of context dfns formals* 49d⟩
 ⟨*Check that all keywords are valid* 65b⟩
 ⟨*Check that namespaces are at the top level* 65c⟩
 ⟨*Code Generator* 25b⟩
 ⟨*codfns.h* 33⟩
 ⟨*Common cell fields* 35⟩
 ⟨*Compile the primitives in prim.apln* 88⟩
 ⟨*Compiler* 23⟩
 ⟨*Compute dfns regions and type, with }* as a child 61c⟩
 ⟨*Compute parser exports* 66b⟩
 ⟨*Compute slots and frames* 62d⟩

⟨Compute the nameclass of *dfns* 61d⟩
 ⟨Compute trad-fns regions 63c⟩
 ⟨Convert ; groups within brackets into *Z* nodes 52c⟩
 ⟨Convert *M* nodes to *F0* nodes 68d⟩
 ⟨Convert α and ω to *V* nodes 49e⟩
 ⟨Convert $\alpha\alpha$ and $\omega\omega$ to *P2* nodes 49f⟩
 ⟨Converters between parent and depth vectors 15c⟩
 ⟨Copy the runtime files into *tests* 89b⟩
 ⟨Count strand and indexing children 50d⟩
 ⟨Declare top-level array structures 50f⟩
 ⟨Declare top-level closures 61b⟩
 ⟨Declare top-level function bindings 61a⟩
 ⟨DWA definitions 43a⟩
 ⟨DWA Function Export 29⟩
 ⟨DWA includes (never defined)⟩
 ⟨DWA macros (never defined)⟩
 ⟨DWA structures and enumerations 42⟩
 ⟨*dwa.c* 47a⟩
 ⟨Enclose $V[X; \dots]$ for expression parsing 53a⟩
 ⟨Global Settings 10a⟩
 ⟨Group function and value expressions 58d⟩
 ⟨Identify label colons vs. others 64d⟩
 ⟨Infer the type of bindings, groups, and variables 54b⟩
 ⟨*init.c* 68a⟩
 ⟨Interface to the backend C compiler 26⟩
 ⟨Lift and flatten expressions 59a⟩
 ⟨Lift *dfns* to the top-level 62a⟩
 ⟨Lift guard tests 64b⟩
 ⟨Line and error reporting utilities 20b⟩
 ⟨Linking with Dyalog 27⟩
 ⟨Map generators over the linearized AST; return 24a⟩
 ⟨Mark APL primitives with appropriate kinds 51c⟩
 ⟨Mark atoms, characters, and numbers as kind 1 50b⟩
 ⟨Mark system variables as *P* nodes with appropriate kinds 52b⟩
 ⟨Must Have APL Utilities 78⟩
 ⟨Nest top-level root lines as *Z* nodes 65d⟩
 ⟨Node \leftrightarrow Generator mapping 24c⟩
 ⟨Node-specific code generators 24e⟩
 ⟨Normalize the input formatting 14b⟩
 ⟨Parse :Namespace syntax 66a⟩
 ⟨Parse assignments 55c⟩
 ⟨Parse Binding nodes 54a⟩
 ⟨Parse brackets and parentheses into $\neg 1$ and *Z* nodes 58c⟩
 ⟨Parse dyadic operator bindings 54c⟩
 ⟨Parse function expressions 60a⟩
 ⟨Parse guards to (*G* (*Z* ...) (*Z* ...)) 64a⟩

⟨*Parse labels* 64f⟩
 ⟨*Parse token stream* 22⟩
 ⟨*Parse trains* 60c⟩
 ⟨*Parse value expressions* 59b⟩
 ⟨*Parser* 17⟩
 ⟨*Parsing Constants* 18a⟩
 ⟨*Prefix code for all generated files* 24d⟩
 ⟨*Pretty-printing AST trees* 79⟩
 ⟨*prim.apln* 32a⟩
 ⟨*Rationalize* $F[X]$ syntax 53e⟩
 ⟨*Rationalize* $V[X; \dots]$ 53b⟩
 ⟨*Record exported top-level bindings* 66c⟩
 ⟨*Set DWA interface functions* 46b⟩
 ⟨*Strand arrays into atoms* 50c⟩
 ⟨*Symbol* \leftrightarrow *Name mapping* 24b⟩
 ⟨*System Primitives* (never defined)⟩
 ⟨*Tangle Commands* 8⟩
 ⟨*The opsys utility* 82b⟩
 ⟨*The Fix API* 13⟩
 ⟨*Tokenize input* 21⟩
 ⟨*Tokenize keywords* 65a⟩
 ⟨*Tokenize labels* 64e⟩
 ⟨*Tokenize numbers* 48b⟩
 ⟨*Tokenize primitives and atoms* 51b⟩
 ⟨*Tokenize strings* 49b⟩
 ⟨*Tokenize system variables* 51d⟩
 ⟨*Tokenize variables* 49c⟩
 ⟨*Unify whitespace and comments* 48a⟩
 ⟨*User-command API* 15a⟩
 ⟨*Verify brackets have function/array target* 52d⟩
 ⟨*Verify source input* ω , set **IN** 14a⟩
 ⟨*Verify that all open characters are valid* 47c⟩
 ⟨*Verify that all structured statements appear within trad-fns* 68c⟩
 ⟨*Verify that system variables are defined* 52a⟩
 ⟨*Wrap all dfns expression bodies as λ nodes* 61e⟩
 ⟨*Wrap expressions as binding or return statements* 62b⟩
 ⟨*XML Rendering* 82a⟩

10.2 Identifiers

AF Δ LIB: 11, 15a, 26, 89a
AF Δ PREFIX: 11, 26
assert: 22, 78
box: 56d, 57a, 58b, 80
box.c: 58a, 58b

cell.c: 41a, 41b
CELL_BOX: 56c, 57a, 57c
cell_type: 35, 36b
CELL_VOID: 36c, 37, 39c
cell_void: 36a, 37, 38a, 38b, 38c, 39a, 40a
codfns: 7, 8, 16b, 24d, 26, 32b, 34a, 41a, 41b, 47b, 58a, 58b, 68a, 68b,
81b, 83, 84b, 84d, 85, 86a, 86c, 86e, 87c, 88, 89a, 89b
codfns.apln: 8
codfns.h: 24d, 33, 34a, 41a, 58a, 68a, 89b
ctyp: 35, 37, 39a, 57a
dct: 79
DECLSPEC: 34b, 37, 38a, 38b, 38c, 39a, 39b, 40a, 40b, 43b, 43c, 44a,
44b, 44c, 44d, 46a, 57a, 57b
DISPLAY: 80, 81a, 81b
DISPLAY.aplf: 81b
dlk: 79
dmx: 20b, 43a, 43b, 44a
dwa.c: 47a, 47b
dwa_dmx: 42, 43a, 44c
dwa_error: 29, 39a, 44a, 44b
dwa_error_ptr: 44a, 44c
dwa_fns: 45, 46a
dwa_wsfns: 45
dwh: 79
dwv: 79
EXPORT: 25a, 34b, 68a
EXPORTING: 34b, 89a
Fix: 13, 15a
init.c: 68a, 68b
lb3: 79
linestarts: 20b
mk_box: 57a, 57b
mk_void: 37, 38a
mkdm: 20b
MKΔRTM: 87b, 87c
MKΔRTM.aplf: 87c
opsys: 26, 82b, 84c, 86d
PP: 27, 81a, 81b
PP.aplf: 81b
pp3: 79
prim: 32a, 32b, 88
prim.apln: 32a, 32b
PS: 13, 17, 88
put: 26, 81c, 87b, 88
quotelines: 20b, 47c, 49a
refc: 35, 37, 38b, 40a, 57a

release_box: 57a, 57b, 57c
release_cell: 39a, 39b, 57a
release_void: 38b, 38c, 39c
retain_cell: 40a, 40b, 50a, 51a, 55b
rgt: 76f, 80
set_codfns_error: 44c, 44d, 46b
set_dmx_message: 43b, 43c
set_dwafns: 24d, 46a
SIGNAL: 14a, 20b, 24a, 24e, 25a, 26, 27, 47c, 48b, 49a, 49d, 52a, 52d,
53e, 54a, 58c, 60a, 61c, 63c, 64a, 64e, 65a, 65b, 65c, 66a, 68c, 78, 81c
src: 8, 13, 16b, 23, 81b, 84d, 86e, 87b, 87c, 88
TANGLE: 83, 84a, 84b, 84c, 84d
TANGLE.aplf: 84d
TANGLE.bat: 84a, 84b
TANGLE.sh: 83, 84a
TEST: 16a, 16b, 64a
TEST.aplf: 16b
tie: 81c, 87b
VERSION: 10b
vsbat: 26, 87b, 89a
vsc: 26, 87b, 89a, 89b
VSΔPATH: 12, 26, 89a
WEAVE: 86a, 86b, 86c, 86d, 86e
WEAVE.aplf: 86e
WEAVE.bat: 86b, 86c
WEAVE.sh: 85, 86a, 86b
wsd: 87b, 89a, 89b
xi: 23, 24a, 25b, 66c
Xml: 82a
xn: 20a, 66b
xt: 20a, 66b
□IO: 10a, 80
□ML: 10a, 80
□WX: 10a

11 GNU AFFERO GPL

Version 3, 19 November 2007
Copyright © 2007 Free Software Foundation, Inc.
<https://fsf.org/>

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

Preamble

The GNU Affero General Public License is a free, copyleft license for software and other kinds of works, specifically designed to ensure cooperation with the community in the case of network server software.

The licenses for most software and other practical works are designed to take away your freedom to share and change the works. By contrast, our General Public Licenses are intended to guarantee your freedom to share and change all versions of a program—to make sure it remains free software for all its users.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for them if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs, and that you know you can do these things. Developers that use our General Public Licenses protect your rights with two steps: (1) assert copyright on the software, and (2) offer you this License which gives you legal permission to copy, distribute and/or modify the software.

A secondary benefit of defending all users' freedom is that improvements made in alternate versions of the program, if they receive widespread use, become available for other developers to incorporate. Many developers of free software are heartened and encouraged by the resulting cooperation. However, in the case of software used on network servers, this result may fail to come about. The GNU General Public License permits making a modified version and letting the public access it on a server without ever releasing its source code to the public.

The GNU Affero General Public License is designed specifically to ensure that, in such cases, the modified source code becomes available to the community. It requires the operator of a network server to provide the source code of the modified version running there to the users of that server. Therefore, public use of a modified version, on a publicly accessible server, gives the public access to the source code of the modified version.

An older license, called the Affero General Public License and published by Affero, was designed to accomplish similar goals. This is a different license, not a version of the Affero GPL, but Affero has released a new version of the Affero GPL which permits relicensing under this license.

The precise terms and conditions for copying, distribution and modification follow.

Terms and Conditions

0. Definitions.

“This License” refers to version 3 of the GNU Affero General Public License.

“Copyright” also means copyright-like laws that apply to other kinds of works, such as semiconductor masks.

“The Program” refers to any copyrightable work licensed under this License. Each licensee is addressed as “you”.

“Licensees” and “recipients” may be individuals or organizations.

To “modify” a work means to copy from or adapt all or part of the work in a fashion requiring copyright permission, other than the making of an exact copy. The resulting work is called a “modified version” of the earlier work or a work “based on” the earlier work.

A “covered work” means either the unmodified Program or a work based on the Program.

To “propagate” a work means to do anything with it that, without permission, would make you directly or secondarily liable for infringement under applicable copyright law, except executing it on a computer or modifying a private copy. Propagation includes copying, distribution (with or without modification), making available to the public, and in some countries other activities as well.

To “convey” a work means any kind of propagation that enables other parties to make or receive copies. Mere interaction with a user through a computer network, with no transfer of a copy, is not conveying.

An interactive user interface displays “Appropriate Legal Notices” to the extent that it includes a convenient and prominently visible feature that (1) displays an appropriate copyright notice, and (2) tells the user that there is no warranty for the work (except to the extent that warranties are provided), that licensees may convey the work under this License, and how to view a copy of this License. If the

interface presents a list of user commands or options, such as a menu, a prominent item in the list meets this criterion.

1. Source Code.

The “source code” for a work means the preferred form of the work for making modifications to it. “Object code” means any non-source form of a work.

A “Standard Interface” means an interface that either is an official standard defined by a recognized standards body, or, in the case of interfaces specified for a particular programming language, one that is widely used among developers working in that language.

The “System Libraries” of an executable work include anything, other than the work as a whole, that (a) is included in the normal form of packaging a Major Component, but which is not part of that Major Component, and (b) serves only to enable use of the work with that Major Component, or to implement a Standard Interface for which an implementation is available to the public in source code form. A “Major Component”, in this context, means a major essential component (kernel, window system, and so on) of the specific operating system (if any) on which the executable work runs, or a compiler used to produce the work, or an object code interpreter used to run it.

The “Corresponding Source” for a work in object code form means all the source code needed to generate, install, and (for an executable work) run the object code and to modify the work, including scripts to control those activities. However, it does not include the work’s System Libraries, or general-purpose tools or generally available free programs which are used unmodified in performing those activities but which are not part of the work. For example, Corresponding Source includes interface definition files associated with source files for the work, and the source code for shared libraries and dynamically linked subprograms that the work is specifically designed to require, such as by intimate data communication or control flow between those subprograms and other parts of the work.

The Corresponding Source need not include anything that users can regenerate automatically from other parts of the Corresponding Source.

The Corresponding Source for a work in source code form is that same work.

2. Basic Permissions.

All rights granted under this License are granted for the term of copyright on the Program, and are irrevocable provided the stated conditions are met. This License explicitly affirms your unlimited permission to run the unmodified Program. The output from running a covered work is covered by this License only if the output, given its content, constitutes a covered work. This License acknowledges your rights of fair use or other equivalent, as provided by copyright law.

You may make, run and propagate covered works that you do not convey, without conditions so long as your license otherwise remains in force. You may convey covered works to others for the sole purpose of having them make modifications exclusively for you, or provide you with facilities for running those works, provided that you comply with the terms of this License in conveying all material for which you do not control copyright. Those thus making or running the covered works for you must do so exclusively on your behalf, under your direction and control, on terms that prohibit them from making any copies of your copyrighted material outside their relationship with you.

Conveying under any other circumstances is permitted solely under the conditions stated below. Sublicensing is not allowed; section 10 makes it unnecessary.

3. Protecting Users' Legal Rights From Anti-Circumvention Law.

No covered work shall be deemed part of an effective technological measure under any applicable law fulfilling obligations under article 11 of the WIPO copyright treaty adopted on 20 December 1996, or similar laws prohibiting or restricting circumvention of such measures.

When you convey a covered work, you waive any legal power to forbid circumvention of technological measures to the extent such circumvention is effected by exercising rights under this License with respect to the covered work, and you disclaim any intention to limit operation or modification of the work as a means of enforcing, against the work's users, your or third parties' legal rights to forbid circumvention of technological measures.

4. Conveying Verbatim Copies.

You may convey verbatim copies of the Program's source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice; keep intact all notices stating that this License and any non-permissive terms added in

accord with section 7 apply to the code; keep intact all notices of the absence of any warranty; and give all recipients a copy of this License along with the Program.

You may charge any price or no price for each copy that you convey, and you may offer support or warranty protection for a fee.

5. Conveying Modified Source Versions.

You may convey a work based on the Program, or the modifications to produce it from the Program, in the form of source code under the terms of section 4, provided that you also meet all of these conditions:

- (a) The work must carry prominent notices stating that you modified it, and giving a relevant date.
- (b) The work must carry prominent notices stating that it is released under this License and any conditions added under section 7. This requirement modifies the requirement in section 4 to “keep intact all notices”.
- (c) You must license the entire work, as a whole, under this License to anyone who comes into possession of a copy. This License will therefore apply, along with any applicable section 7 additional terms, to the whole of the work, and all its parts, regardless of how they are packaged. This License gives no permission to license the work in any other way, but it does not invalidate such permission if you have separately received it.
- (d) If the work has interactive user interfaces, each must display Appropriate Legal Notices; however, if the Program has interactive interfaces that do not display Appropriate Legal Notices, your work need not make them do so.

A compilation of a covered work with other separate and independent works, which are not by their nature extensions of the covered work, and which are not combined with it such as to form a larger program, in or on a volume of a storage or distribution medium, is called an “aggregate” if the compilation and its resulting copyright are not used to limit the access or legal rights of the compilation’s users beyond what the individual works permit. Inclusion of a covered work in an aggregate does not cause this License to apply to the other parts of the aggregate.

6. Conveying Non-Source Forms.

You may convey a covered work in object code form under the terms of sections 4 and 5, provided that you also convey the machine-readable Corresponding Source under the terms of this License, in one of these ways:

- (a) Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by the Corresponding Source fixed on a durable physical medium customarily used for software interchange.
- (b) Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by a written offer, valid for at least three years and valid for as long as you offer spare parts or customer support for that product model, to give anyone who possesses the object code either (1) a copy of the Corresponding Source for all the software in the product that is covered by this License, on a durable physical medium customarily used for software interchange, for a price no more than your reasonable cost of physically performing this conveying of source, or (2) access to copy the Corresponding Source from a network server at no charge.
- (c) Convey individual copies of the object code with a copy of the written offer to provide the Corresponding Source. This alternative is allowed only occasionally and noncommercially, and only if you received the object code with such an offer, in accord with subsection 6b.
- (d) Convey the object code by offering access from a designated place (gratis or for a charge), and offer equivalent access to the Corresponding Source in the same way through the same place at no further charge. You need not require recipients to copy the Corresponding Source along with the object code. If the place to copy the object code is a network server, the Corresponding Source may be on a different server (operated by you or a third party) that supports equivalent copying facilities, provided you maintain clear directions next to the object code saying where to find the Corresponding Source. Regardless of what server hosts the Corresponding Source, you remain obligated to ensure that it is available for as long as needed to satisfy these requirements.
- (e) Convey the object code using peer-to-peer transmission, provided you inform other peers where the object code

and Corresponding Source of the work are being offered to the general public at no charge under subsection 6d.

A separable portion of the object code, whose source code is excluded from the Corresponding Source as a System Library, need not be included in conveying the object code work.

A “User Product” is either (1) a “consumer product”, which means any tangible personal property which is normally used for personal, family, or household purposes, or (2) anything designed or sold for incorporation into a dwelling. In determining whether a product is a consumer product, doubtful cases shall be resolved in favor of coverage. For a particular product received by a particular user, “normally used” refers to a typical or common use of that class of product, regardless of the status of the particular user or of the way in which the particular user actually uses, or expects or is expected to use, the product. A product is a consumer product regardless of whether the product has substantial commercial, industrial or non-consumer uses, unless such uses represent the only significant mode of use of the product.

“Installation Information” for a User Product means any methods, procedures, authorization keys, or other information required to install and execute modified versions of a covered work in that User Product from a modified version of its Corresponding Source. The information must suffice to ensure that the continued functioning of the modified object code is in no case prevented or interfered with solely because modification has been made.

If you convey an object code work under this section in, or with, or specifically for use in, a User Product, and the conveying occurs as part of a transaction in which the right of possession and use of the User Product is transferred to the recipient in perpetuity or for a fixed term (regardless of how the transaction is characterized), the Corresponding Source conveyed under this section must be accompanied by the Installation Information. But this requirement does not apply if neither you nor any third party retains the ability to install modified object code on the User Product (for example, the work has been installed in ROM).

The requirement to provide Installation Information does not include a requirement to continue to provide support service, warranty, or updates for a work that has been modified or installed by the recipient, or for the User Product in which it has been modified or installed. Access to a network may be denied when the modification itself materially and adversely

affects the operation of the network or violates the rules and protocols for communication across the network.

Corresponding Source conveyed, and Installation Information provided, in accord with this section must be in a format that is publicly documented (and with an implementation available to the public in source code form), and must require no special password or key for unpacking, reading or copying.

7. Additional Terms.

“Additional permissions” are terms that supplement the terms of this License by making exceptions from one or more of its conditions. Additional permissions that are applicable to the entire Program shall be treated as though they were included in this License, to the extent that they are valid under applicable law. If additional permissions apply only to part of the Program, that part may be used separately under those permissions, but the entire Program remains governed by this License without regard to the additional permissions.

When you convey a copy of a covered work, you may at your option remove any additional permissions from that copy, or from any part of it. (Additional permissions may be written to require their own removal in certain cases when you modify the work.) You may place additional permissions on material, added by you to a covered work, for which you have or can give appropriate copyright permission.

Notwithstanding any other provision of this License, for material you add to a covered work, you may (if authorized by the copyright holders of that material) supplement the terms of this License with terms:

- (a) Disclaiming warranty or limiting liability differently from the terms of sections 15 and 16 of this License; or
- (b) Requiring preservation of specified reasonable legal notices or author attributions in that material or in the Appropriate Legal Notices displayed by works containing it; or
- (c) Prohibiting misrepresentation of the origin of that material, or requiring that modified versions of such material be marked in reasonable ways as different from the original version; or
- (d) Limiting the use for publicity purposes of names of licensors or authors of the material; or
- (e) Declining to grant rights under trademark law for use of some trade names, trademarks, or service marks; or

- (f) Requiring indemnification of licensors and authors of that material by anyone who conveys the material (or modified versions of it) with contractual assumptions of liability to the recipient, for any liability that these contractual assumptions directly impose on those licensors and authors.

All other non-permissive additional terms are considered “further restrictions” within the meaning of section 10. If the Program as you received it, or any part of it, contains a notice stating that it is governed by this License along with a term that is a further restriction, you may remove that term. If a license document contains a further restriction but permits relicensing or conveying under this License, you may add to a covered work material governed by the terms of that license document, provided that the further restriction does not survive such relicensing or conveying.

If you add terms to a covered work in accord with this section, you must place, in the relevant source files, a statement of the additional terms that apply to those files, or a notice indicating where to find the applicable terms.

Additional terms, permissive or non-permissive, may be stated in the form of a separately written license, or stated as exceptions; the above requirements apply either way.

8. Termination.

You may not propagate or modify a covered work except as expressly provided under this License. Any attempt otherwise to propagate or modify it is void, and will automatically terminate your rights under this License (including any patent licenses granted under the third paragraph of section 11).

However, if you cease all violation of this License, then your license from a particular copyright holder is reinstated (a) provisionally, unless and until the copyright holder explicitly and finally terminates your license, and (b) permanently, if the copyright holder fails to notify you of the violation by some reasonable means prior to 60 days after the cessation.

Moreover, your license from a particular copyright holder is reinstated permanently if the copyright holder notifies you of the violation by some reasonable means, this is the first time you have received notice of violation of this License (for any work) from that copyright holder, and you cure the violation prior to 30 days after your receipt of the notice.

Termination of your rights under this section does not terminate the licenses of parties who have received copies or

rights from you under this License. If your rights have been terminated and not permanently reinstated, you do not qualify to receive new licenses for the same material under section 10.

9. Acceptance Not Required for Having Copies.

You are not required to accept this License in order to receive or run a copy of the Program. Ancillary propagation of a covered work occurring solely as a consequence of using peer-to-peer transmission to receive a copy likewise does not require acceptance. However, nothing other than this License grants you permission to propagate or modify any covered work. These actions infringe copyright if you do not accept this License. Therefore, by modifying or propagating a covered work, you indicate your acceptance of this License to do so.

10. Automatic Licensing of Downstream Recipients.

Each time you convey a covered work, the recipient automatically receives a license from the original licensors, to run, modify and propagate that work, subject to this License. You are not responsible for enforcing compliance by third parties with this License.

An “entity transaction” is a transaction transferring control of an organization, or substantially all assets of one, or subdividing an organization, or merging organizations. If propagation of a covered work results from an entity transaction, each party to that transaction who receives a copy of the work also receives whatever licenses to the work the party’s predecessor in interest had or could give under the previous paragraph, plus a right to possession of the Corresponding Source of the work from the predecessor in interest, if the predecessor has it or can get it with reasonable efforts.

You may not impose any further restrictions on the exercise of the rights granted or affirmed under this License. For example, you may not impose a license fee, royalty, or other charge for exercise of rights granted under this License, and you may not initiate litigation (including a cross-claim or counterclaim in a lawsuit) alleging that any patent claim is infringed by making, using, selling, offering for sale, or importing the Program or any portion of it.

11. Patents.

A “contributor” is a copyright holder who authorizes use under this License of the Program or a work on which the

Program is based. The work thus licensed is called the contributor's "contributor version".

A contributor's "essential patent claims" are all patent claims owned or controlled by the contributor, whether already acquired or hereafter acquired, that would be infringed by some manner, permitted by this License, of making, using, or selling its contributor version, but do not include claims that would be infringed only as a consequence of further modification of the contributor version. For purposes of this definition, "control" includes the right to grant patent sublicenses in a manner consistent with the requirements of this License.

Each contributor grants you a non-exclusive, worldwide, royalty-free patent license under the contributor's essential patent claims, to make, use, sell, offer for sale, import and otherwise run, modify and propagate the contents of its contributor version.

In the following three paragraphs, a "patent license" is any express agreement or commitment, however denominated, not to enforce a patent (such as an express permission to practice a patent or covenant not to sue for patent infringement). To "grant" such a patent license to a party means to make such an agreement or commitment not to enforce a patent against the party.

If you convey a covered work, knowingly relying on a patent license, and the Corresponding Source of the work is not available for anyone to copy, free of charge and under the terms of this License, through a publicly available network server or other readily accessible means, then you must either (1) cause the Corresponding Source to be so available, or (2) arrange to deprive yourself of the benefit of the patent license for this particular work, or (3) arrange, in a manner consistent with the requirements of this License, to extend the patent license to downstream recipients. "Knowingly relying" means you have actual knowledge that, but for the patent license, your conveying the covered work in a country, or your recipient's use of the covered work in a country, would infringe one or more identifiable patents in that country that you have reason to believe are valid.

If, pursuant to or in connection with a single transaction or arrangement, you convey, or propagate by procuring conveyance of, a covered work, and grant a patent license to some of the parties receiving the covered work authorizing them to use, propagate, modify or convey a specific copy of

the covered work, then the patent license you grant is automatically extended to all recipients of the covered work and works based on it.

A patent license is “discriminatory” if it does not include within the scope of its coverage, prohibits the exercise of, or is conditioned on the non-exercise of one or more of the rights that are specifically granted under this License. You may not convey a covered work if you are a party to an arrangement with a third party that is in the business of distributing software, under which you make payment to the third party based on the extent of your activity of conveying the work, and under which the third party grants, to any of the parties who would receive the covered work from you, a discriminatory patent license (a) in connection with copies of the covered work conveyed by you (or copies made from those copies), or (b) primarily for and in connection with specific products or compilations that contain the covered work, unless you entered into that arrangement, or that patent license was granted, prior to 28 March 2007.

Nothing in this License shall be construed as excluding or limiting any implied license or other defenses to infringement that may otherwise be available to you under applicable patent law.

12. No Surrender of Others’ Freedom.

If conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot convey a covered work so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not convey it at all. For example, if you agree to terms that obligate you to collect a royalty for further conveying from those to whom you convey the Program, the only way you could satisfy both those terms and this License would be to refrain entirely from conveying the Program.

13. Remote Network Interaction; Use with the GNU General Public License.

Notwithstanding any other provision of this License, if you modify the Program, your modified version must prominently offer all users interacting with it remotely through a computer network (if your version supports such interaction) an opportunity to receive the Corresponding Source of your version by providing access to the

Corresponding Source from a network server at no charge, through some standard or customary means of facilitating copying of software. This Corresponding Source shall include the Corresponding Source for any work covered by version 3 of the GNU General Public License that is incorporated pursuant to the following paragraph.

Notwithstanding any other provision of this License, you have permission to link or combine any covered work with a work licensed under version 3 of the GNU General Public License into a single combined work, and to convey the resulting work. The terms of this License will continue to apply to the part which is the covered work, but the work with which it is combined will remain governed by version 3 of the GNU General Public License.

14. Revised Versions of this License.

The Free Software Foundation may publish revised and/or new versions of the GNU Affero General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies that a certain numbered version of the GNU Affero General Public License “or any later version” applies to it, you have the option of following the terms and conditions either of that numbered version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of the GNU Affero General Public License, you may choose any version ever published by the Free Software Foundation.

If the Program specifies that a proxy can decide which future versions of the GNU Affero General Public License can be used, that proxy’s public statement of acceptance of a version permanently authorizes you to choose that version for the Program.

Later license versions may give you additional or different permissions. However, no additional obligations are imposed on any author or copyright holder as a result of your choosing to follow a later version.

15. Disclaimer of Warranty.

THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES

PROVIDE THE PROGRAM “AS IS” WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.

16. Limitation of Liability.

IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MODIFIES AND/OR CONVEYS THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

17. Interpretation of Sections 15 and 16.

If the disclaimer of warranty and limitation of liability provided above cannot be given local legal effect according to their terms, reviewing courts shall apply local law that most closely approximates an absolute waiver of all civil liability in connection with the Program, unless a warranty or assumption of liability accompanies a copy of the Program in return for a fee.

End of Terms and Conditions

How to Apply These Terms to Your New Programs

If you develop a new program, and you want it to be of the greatest possible use to the public, the best way to achieve this is to make it free software which everyone can redistribute and change under these terms.

To do so, attach the following notices to the program. It is safest to attach them to the start of each source file to most

effectively state the exclusion of warranty; and each file should have at least the “copyright” line and a pointer to where the full notice is found.

```
<one line to give the program's name and a brief idea of what it does.>
```

```
Copyright (C) <textyear> <name of author>
```

```
This program is free software: you can redistribute it and/or  
modify it under the terms of the GNU Affero General Public  
License as published by the Free Software Foundation, either  
version 3 of the License, or (at your option) any later version.
```

```
This program is distributed in the hope that it will be useful,  
but WITHOUT ANY WARRANTY; without even the implied warranty of  
MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU  
Affero General Public License for more details.
```

```
You should have received a copy of the GNU Affero General Public  
License along with this program. If not, see  
<https://www.gnu.org/licenses/>.
```

Also add information on how to contact you by electronic and paper mail.

If your software can interact with users remotely through a computer network, you should also make sure that it provides a way for users to get its source. For example, if your program is a web application, its interface could display a “Source” link that leads users to an archive of the code. There are many ways you could offer source, and different solutions will be better for different programs; see section 13 for the specific requirements.

You should also get your employer (if you work as a programmer) or school, if any, to sign a “copyright disclaimer” for the program, if necessary. For more information on this, and how to apply and follow the GNU AGPL, see <https://www.gnu.org/licenses/>.